

Convolutional NNs and Generative Models

Friday
08h00-09h00

TO COMPLETE YOUR REGISTRATION, PLEASE TELL US WHETHER OR NOT THIS IMAGE CONTAINS A STOP SIGN:



NO

YES

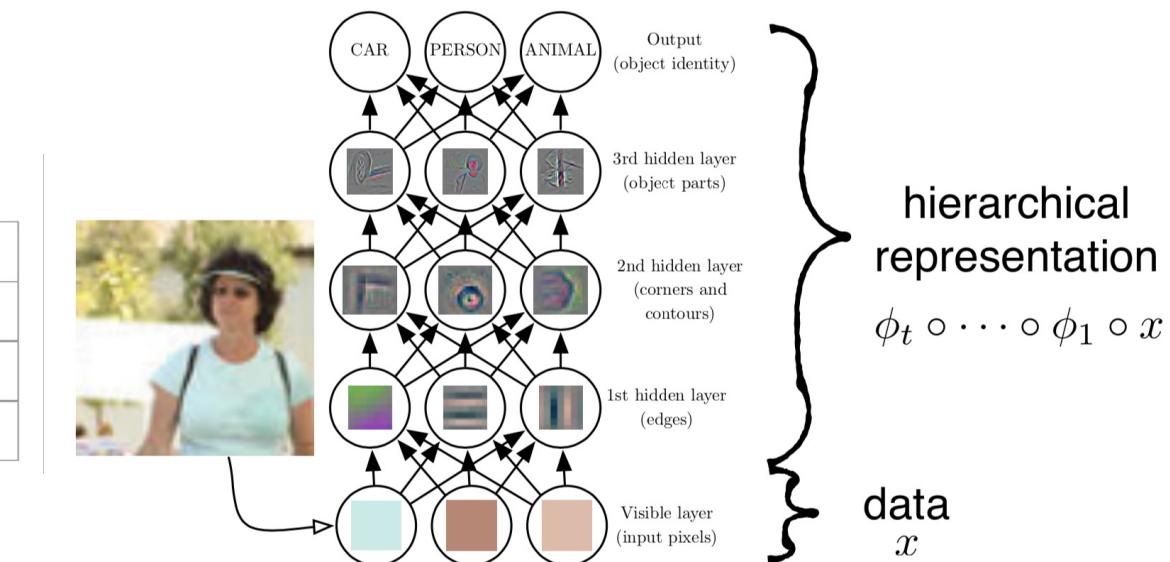
ANSWER QUICKLY—OUR SELF-DRIVING CAR IS ALMOST AT THE INTERSECTION.

SO MUCH OF "AI" IS JUST FIGURING OUT WAYS TO OFFLOAD WORK ONTO RANDOM STRANGERS.

Convolutional Neural Networks

- Specialized Neural Network for data arranged **on a grid**
 - Images
 - DNA sequences
 - ...

	A	C	G	T	W	S	M	K	R	Y	B	D	H	V	N	Z
A	1	0	0	0	1/2	0	1/2	0	1/2	0	0	1/3	1/3	1/3	1/4	0
C	0	1	0	0	0	1/2	1/2	0	0	1/2	1/3	0	1/3	1/3	1/4	0
G	0	0	1	0	0	1/2	0	1/2	1/2	0	1/3	1/3	0	1/3	1/4	0
T	0	0	0	1	1/2	0	0	1/2	0	1/2	1/3	1/3	1/3	0	1/4	0



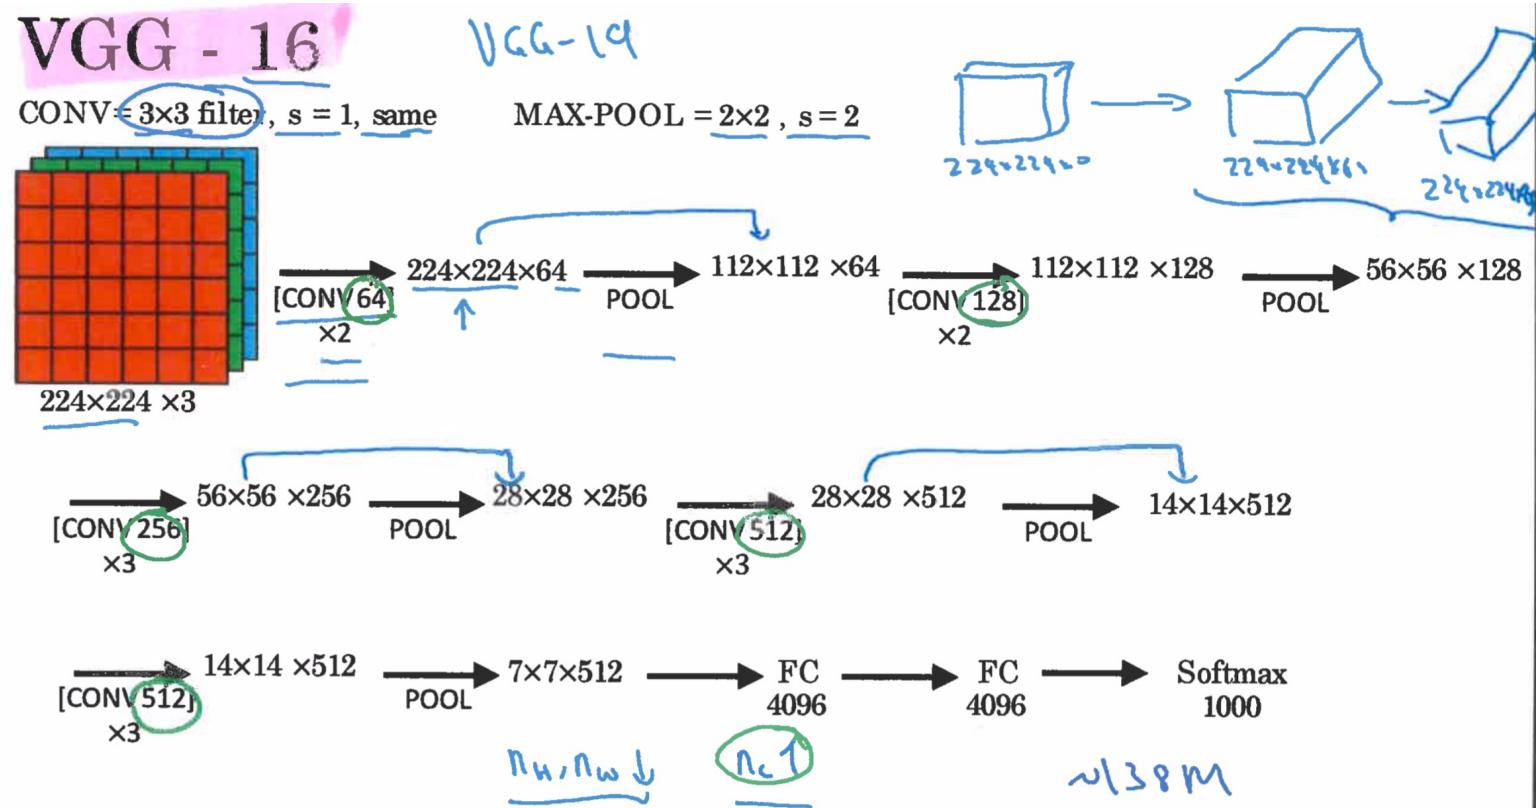
Types of layer in CNN

- Convolution (CONV)

- Pooling (POOL)

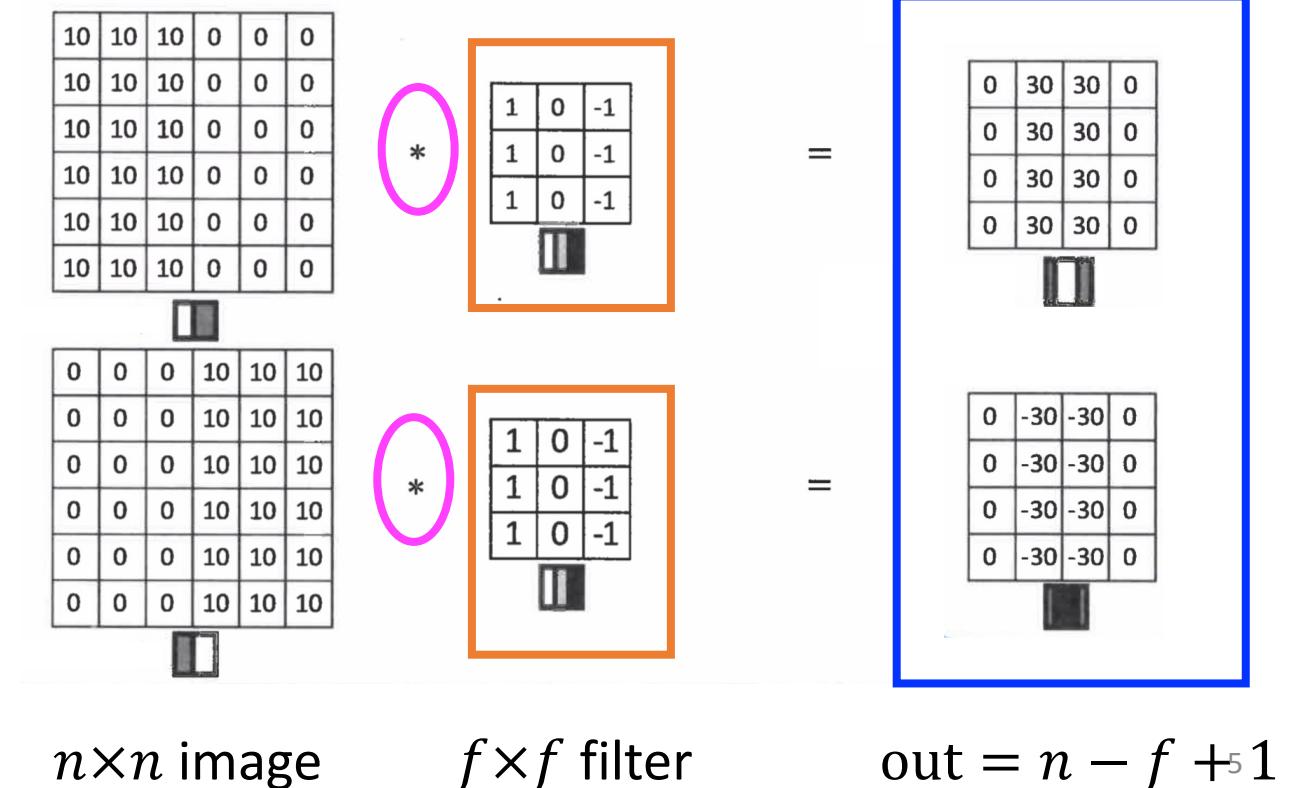
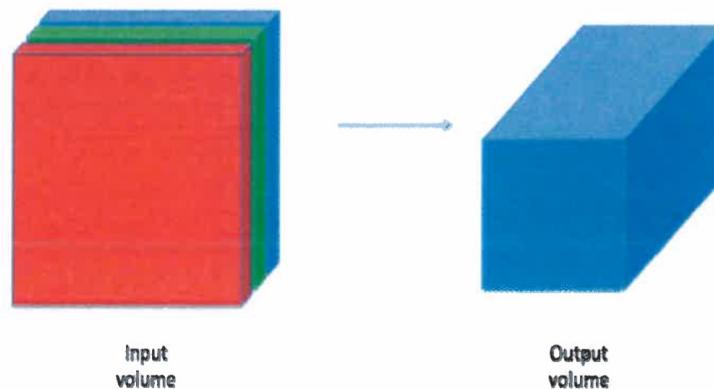
- Fully connected (FC)

• *Usually multiple CONV layers followed by a POOL layer, and FC layers in the last few layers*

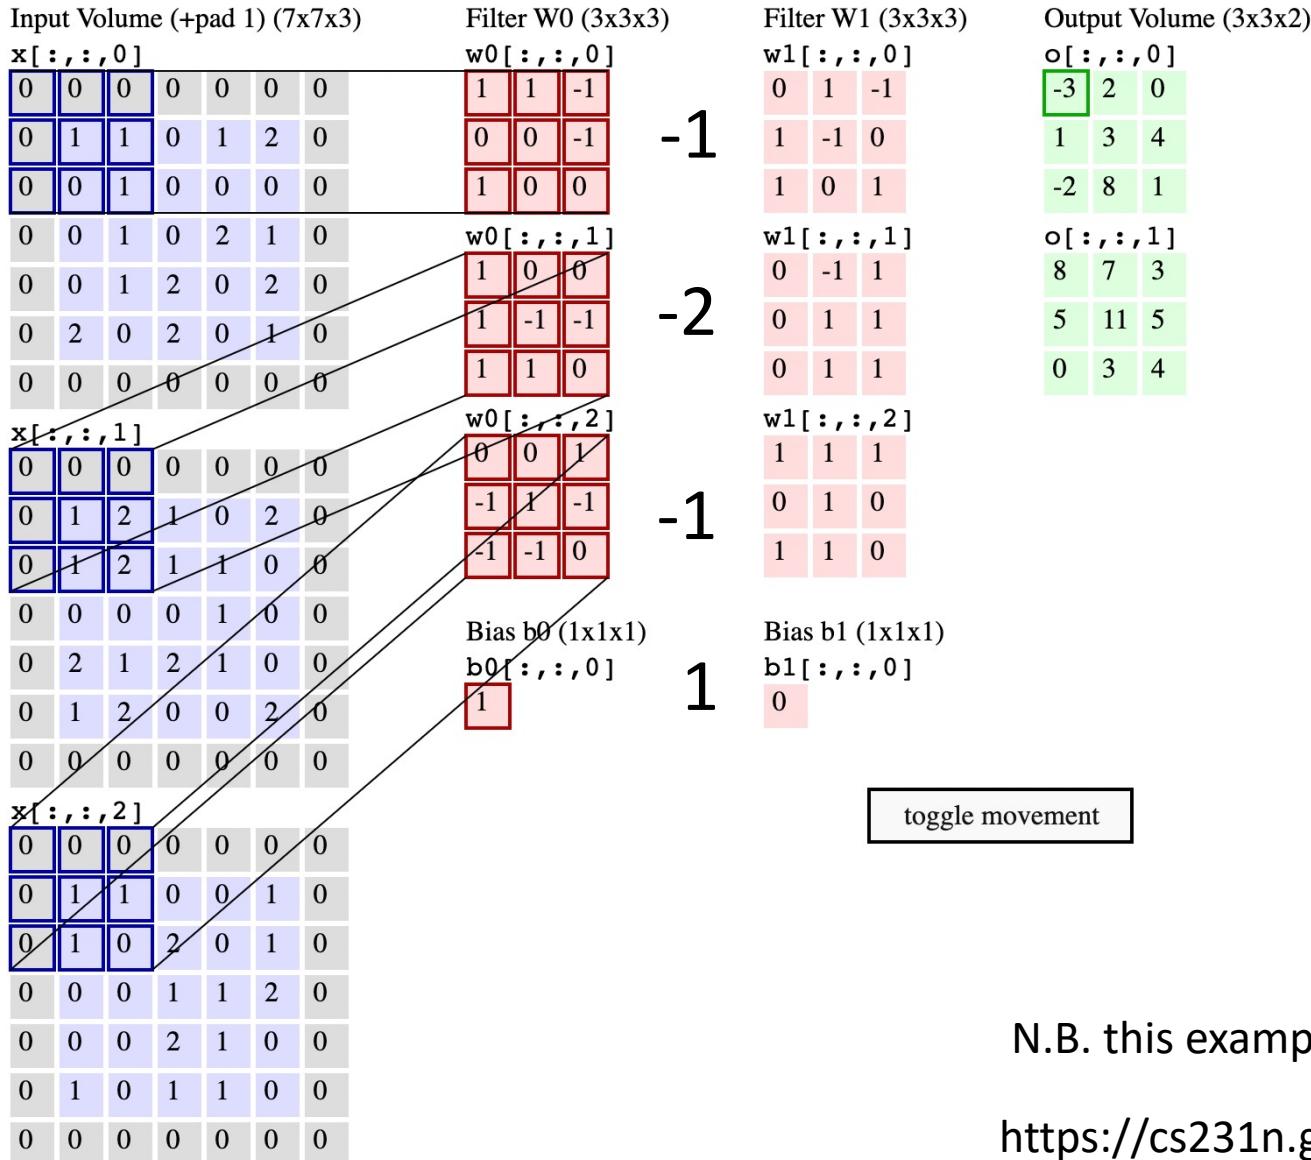


Convolution Layer (CONV)

- **Convolution** transforms an input volume into an **output volume** of different size, also called **feature map**
- **Filter kernels** are used to detect features (for example, edge detection in 1st hidden layer)



Convolution Layer (CONV)

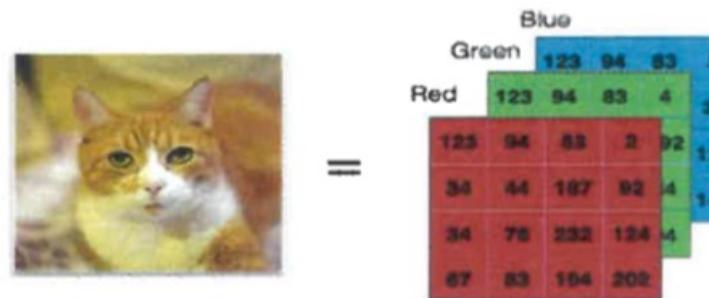


N.B. this example the **stride = 2**, **padding = 1**, **channels = 2**

<https://cs231n.github.io/convolutional-networks/>

Padding

Adds zeros around the border of an image



pad

pad

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	123 94 83 2	0	0	0	0	0	0
0	0	34 44 187 92	0	0	0	0	0	0
0	0	34 78 232 124	0	0	0	0	0	0
0	0	67 83 154 202	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

pad

pad

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	123 94 2 4	0	0	0	0	0	0
0	0	11 3 22 192	0	0	0	0	0	0
0	0	12 4 23 34	0	0	0	0	0	0
0	0	194 83 12 94	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Use cases :

- Keeps more information at the border of an image
- Allows to *use a CONV layer without shrinking* the height and width of the volumes (important for deeper networks)

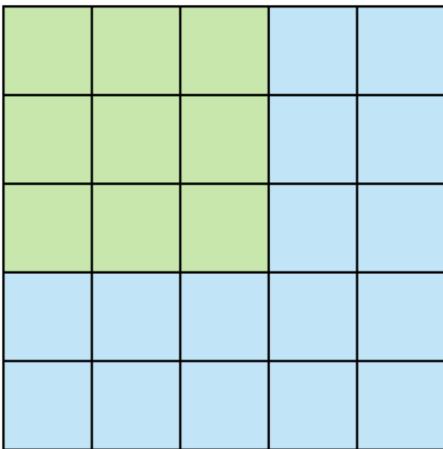
pad

pad

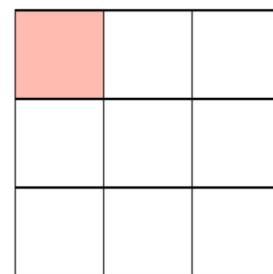
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	123 94 83 2	0	0	0	0	0	0
0	0	34 44 37 30	0	0	0	0	0	0
0	0	34 114 234 134	0	0	0	0	0	0
0	0	49 18 204 143	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

Strided convolutions

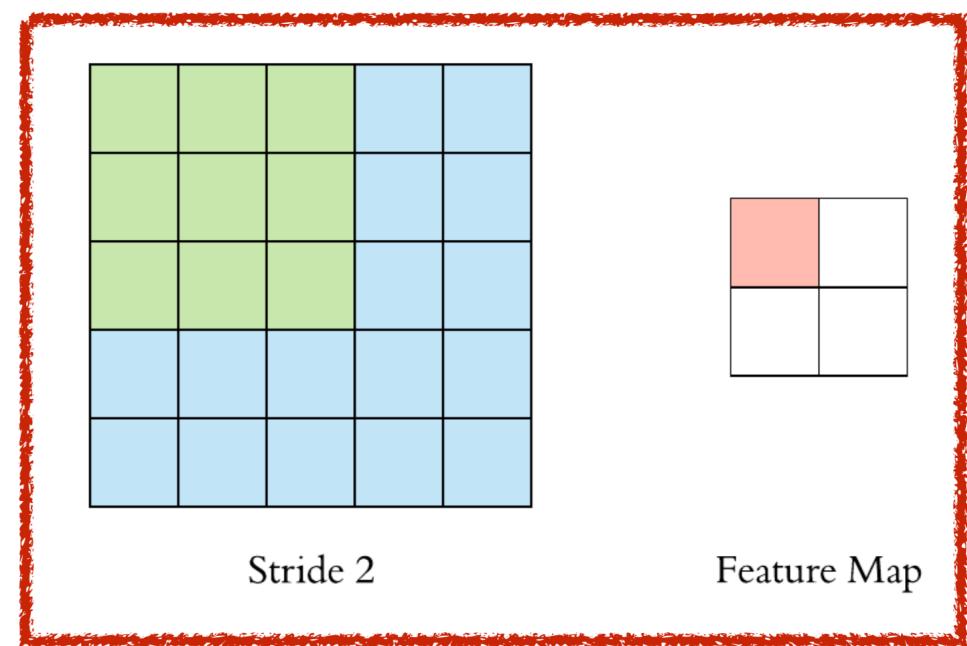
By how much you move the filter



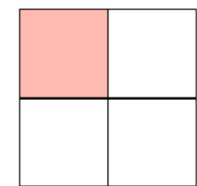
Stride 1



Feature Map



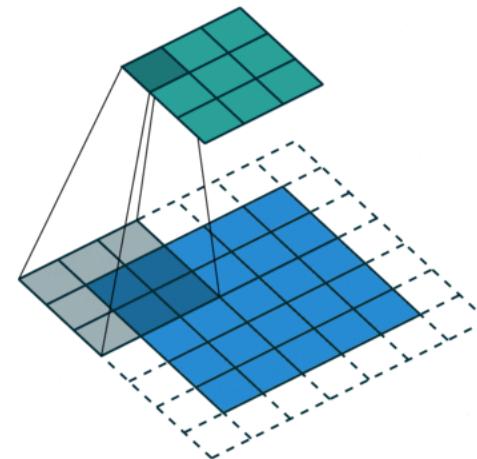
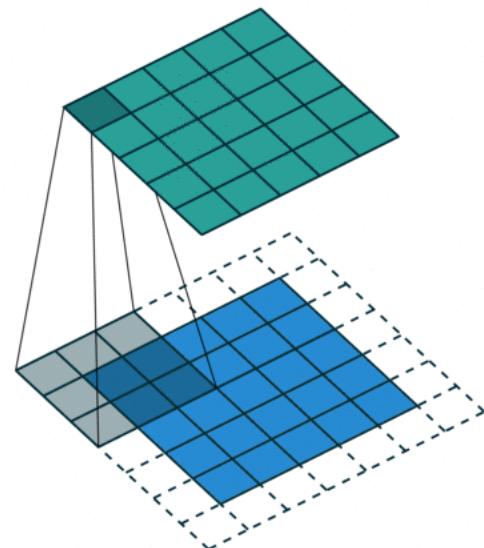
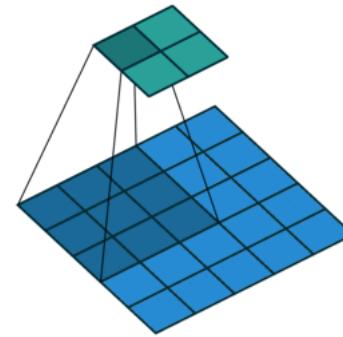
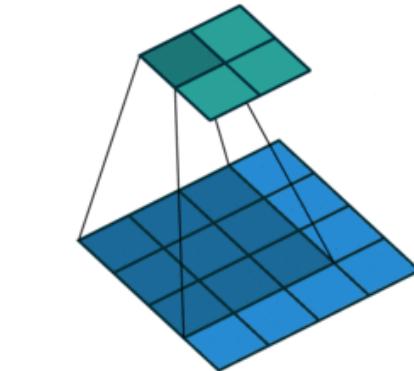
Stride 2



Feature Map

increasing stride from 1 to 2

Illustration



Pooling Layer (POOL)

- reduces the spatial dimension (n_H and n_W) to **decrease computational power**

Max Pool

2	3	1	9
4	7	3	5
8	2	2	2
1	3	4	5



7	9
8	5

Max-Pool with a
2 by 2 filter and
stride 2.

Get the max value

Average Pool

2	3	1	9
4	7	3	5
8	2	2	2
1	3	4	5



4	4.5
3.25	3.25

Average Pool with
a 2 by 2 filter and
stride 2.

Get the average value

One Look Is Worth A Thousand Words--

One look at our line of Republic, Firestone, Miller and United States tires can tell you more than a hundred personal letters or advertisements.

**WE WILL PROVE THEIR VALUE
BEFORE YOU INVEST ONE DOLLAR
IN THEM.**

Ever consider buying Supplies from a catalog?

What's the use! Call and see what you are buying. One look at our display of automobile and motorcycle accessories will convince you of the fact.

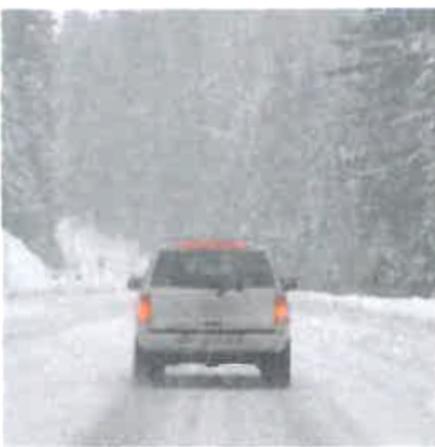
**THAT WE HAVE EVERYTHING FOR
THE AUTO**

Piqua Auto Supply House

133 N. Main St.—Piqua, O.

Object Detection

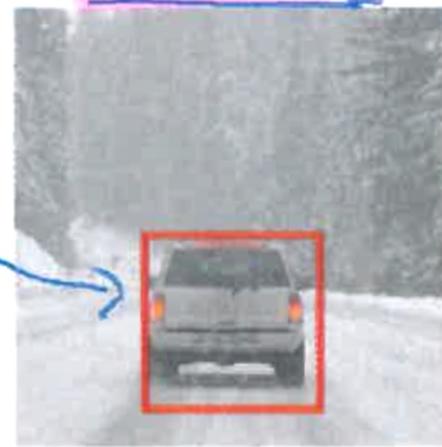
Image classification



"Car"

1 object

Classification with localization



"Car"

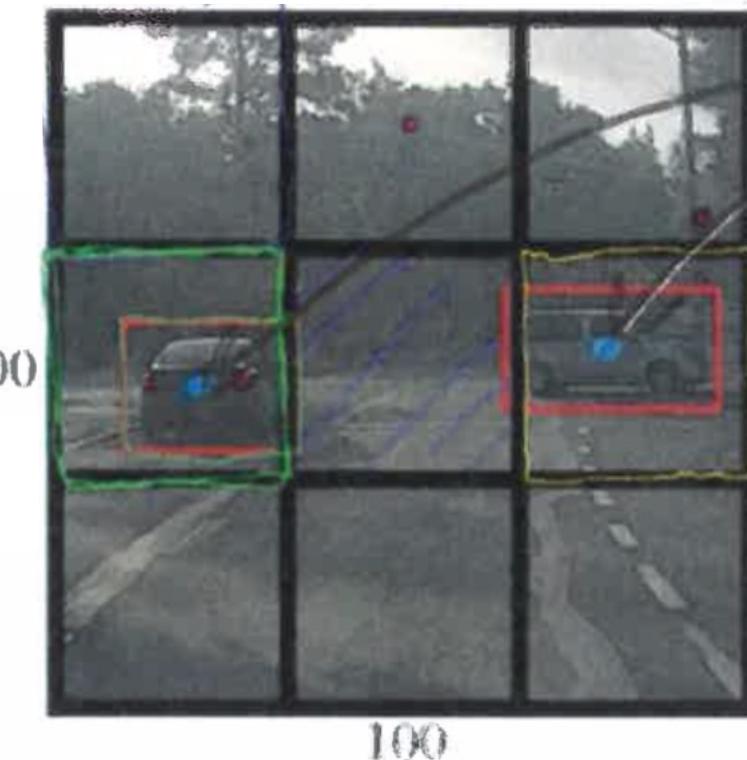
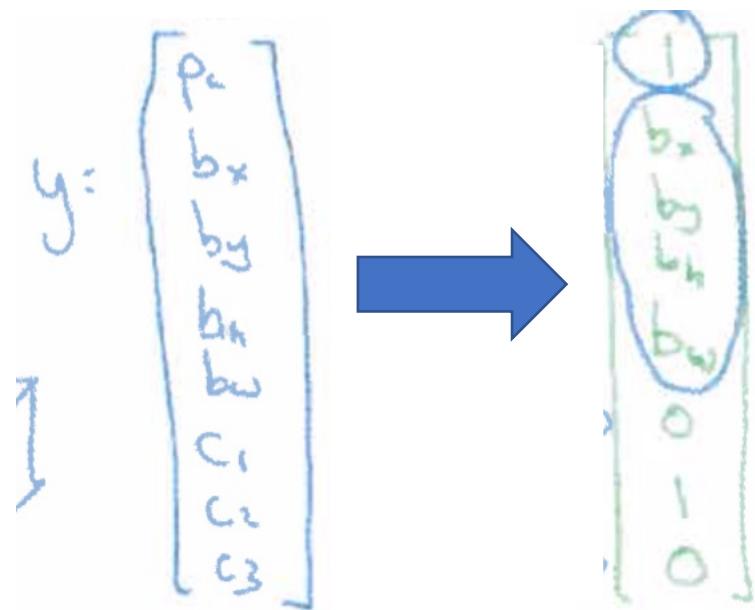
Detection



multiple
objects

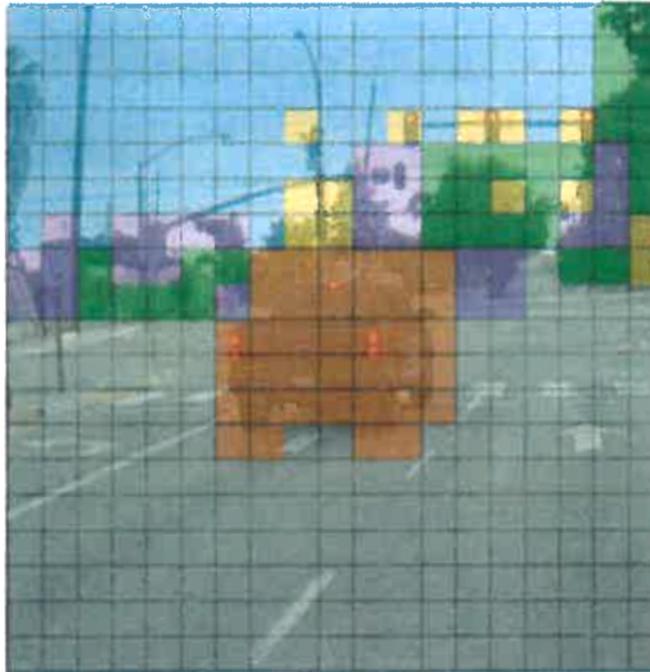
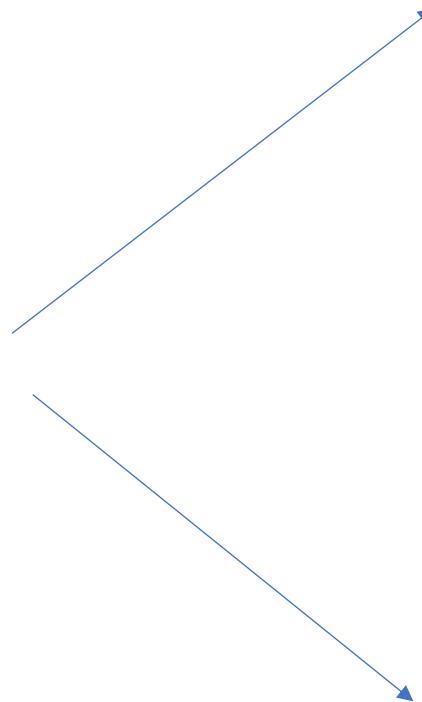
Yolo (YOu Look only Once)

- define a grid in the image
- apply the training to each cell (need ground-truth bounding boxes)
- For each 'anchor box' and 3 classes we have:



- Allows for overlapping objects

YOLO prediction visualisation

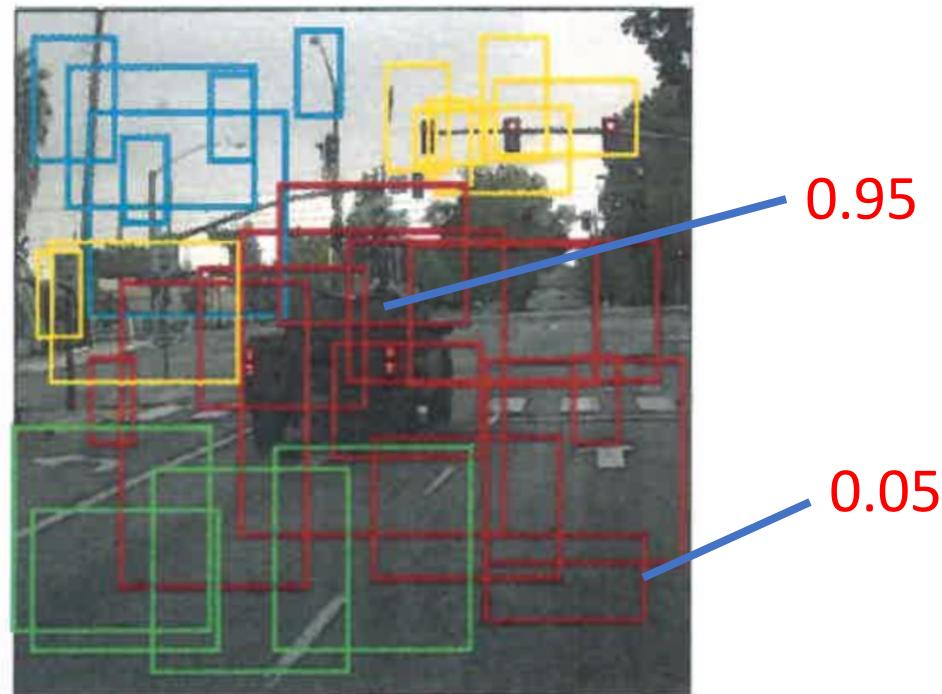


car
road sign
tree
traffic light
sky
background

- Filter the boxes using :
 - 1) score thresholding
 - 2) non-max suppression

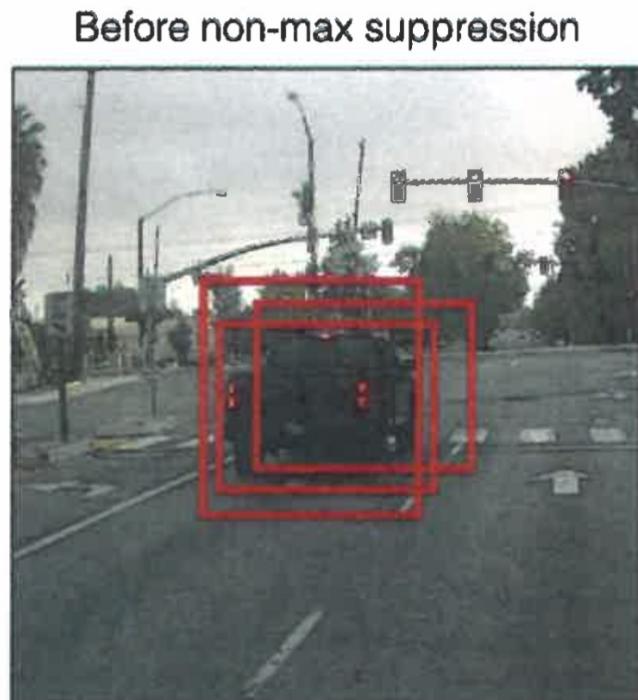
Score Thresholding

- Throw away boxes that have detected a class with a score less than the threshold (*0.6 for example*)



Non-max suppression

- ensures that an object is detected **only once**
 - Choice based on the p_c value : *keep the largest p_c output and discard any remaining box with $IoU > 0.5$*



Non-Max
Suppression



Intersection Over Union (IOU)

- performance metric on how similar two boxes are with each other
- (higher the better!)

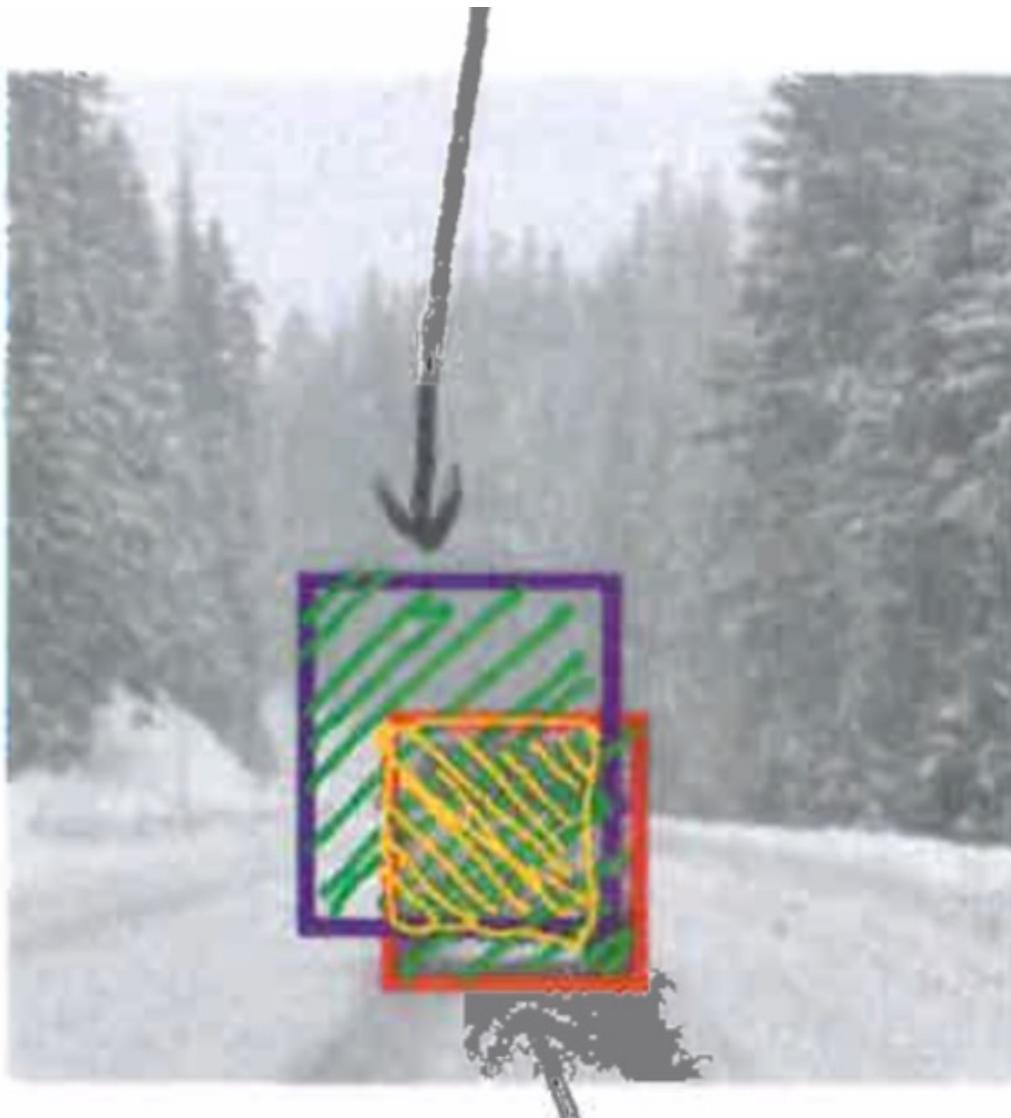
Outcome of the algorithm



True bounding box

$$\text{IoU} = \frac{\text{Size of intersection}}{\text{Size of union}}$$

Intersection Over Union (IOU)



Yellow = intersection I

Green = union U

$$\text{IoU} = I/U$$

Can express this as:

$$TP / (TP + FN + FP)$$

[see also, Jaccard Index]

Generative Models

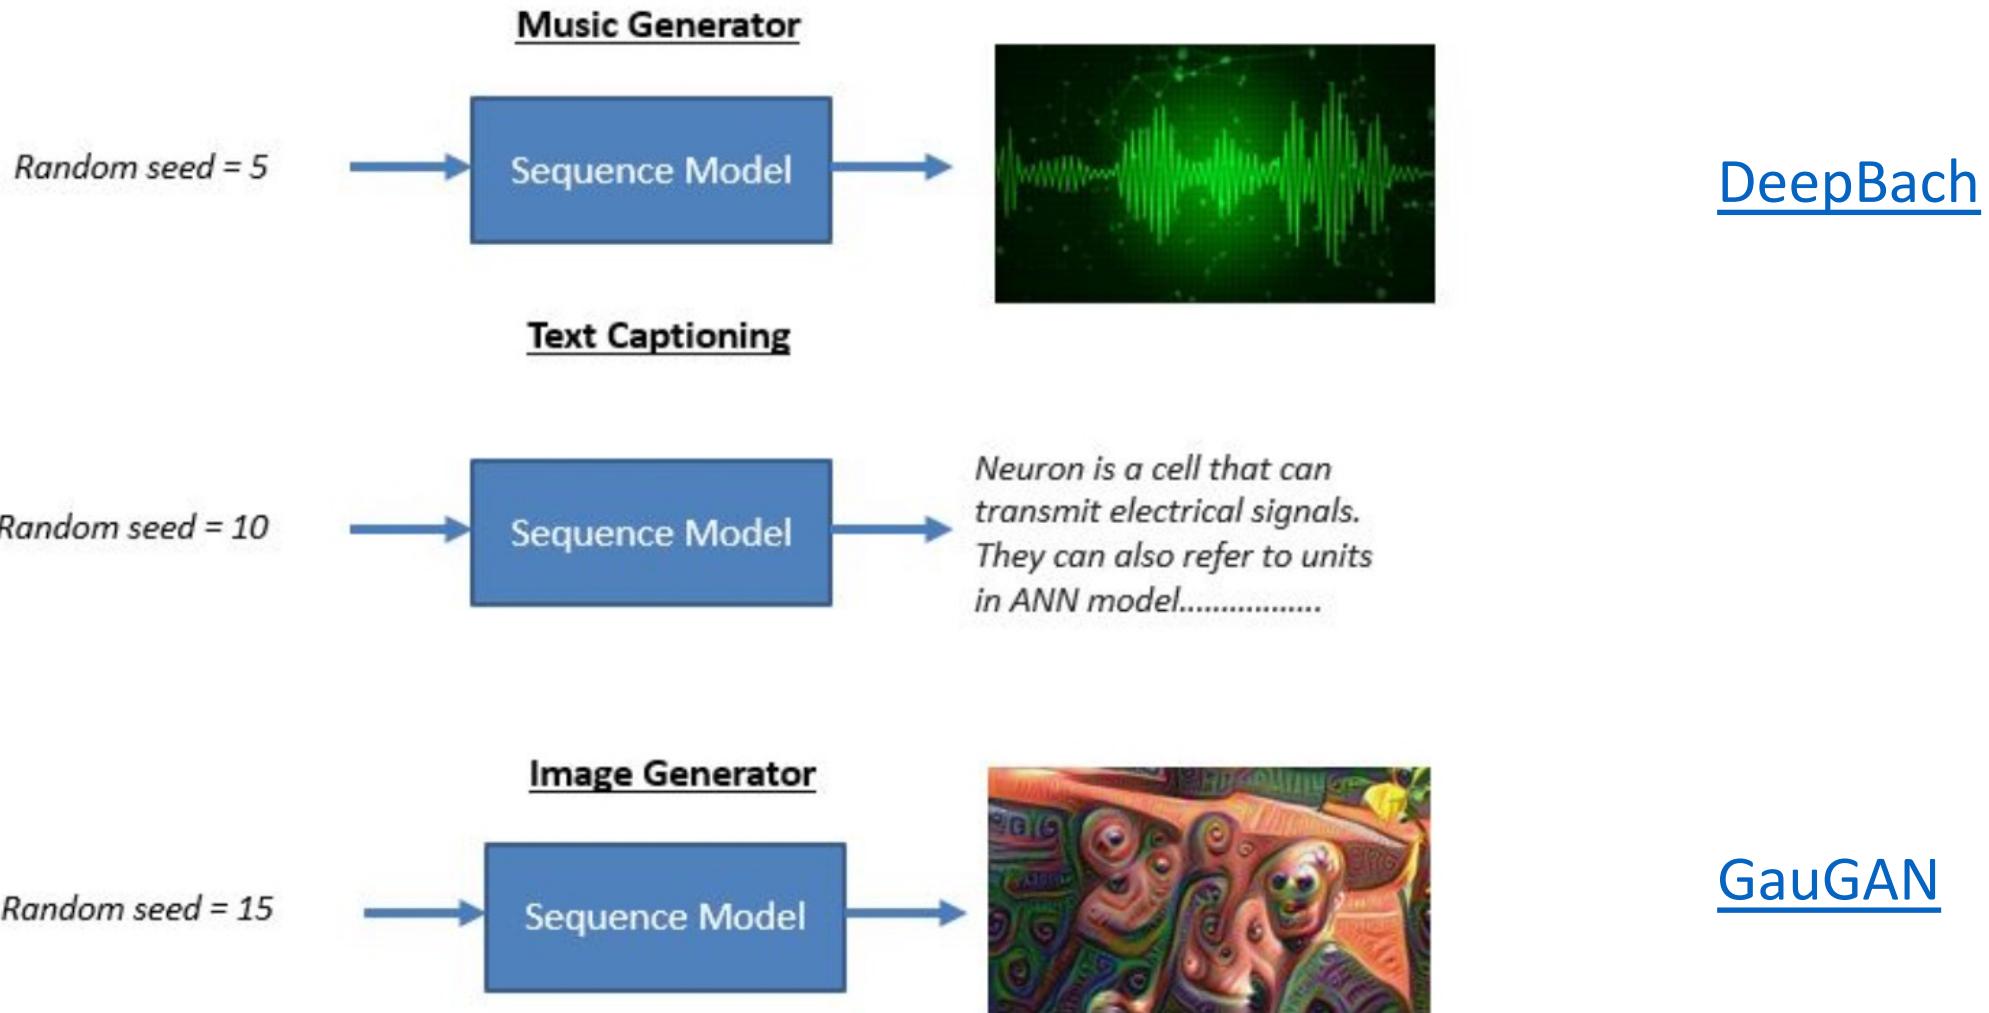
Holy grail of Deep Learning
these days



DALL-E : 'A photograph of a cow on the moon.'



Generative Models Examples

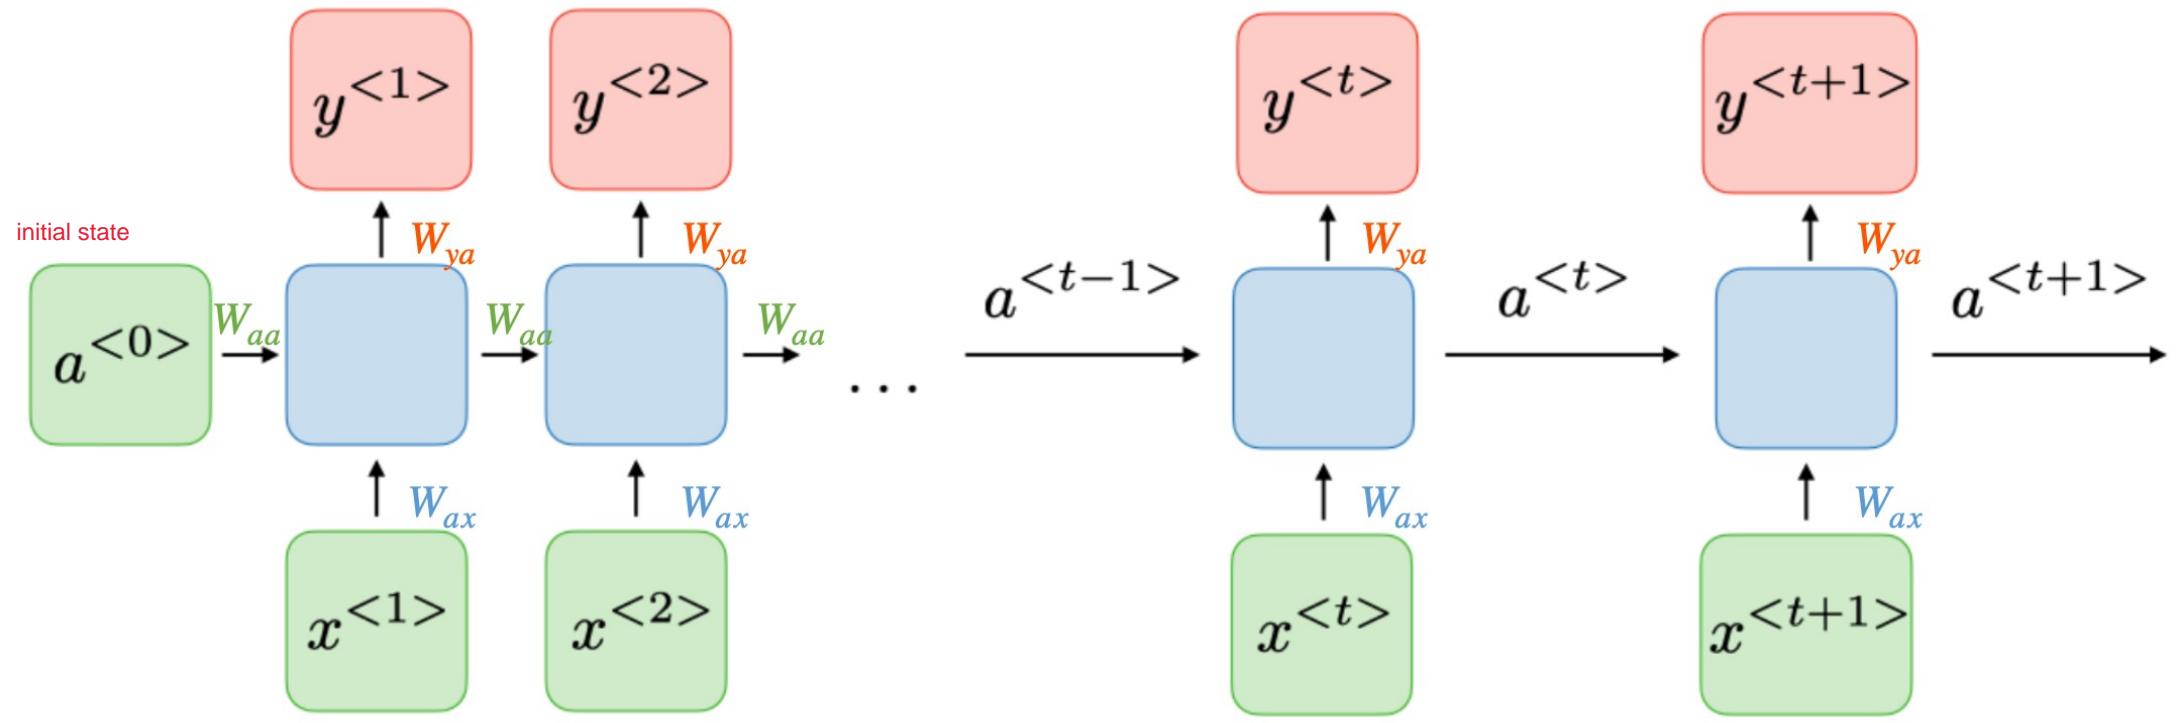


[DeepBach](#)

[GauGAN](#)

RNN model

recurrent neural networks



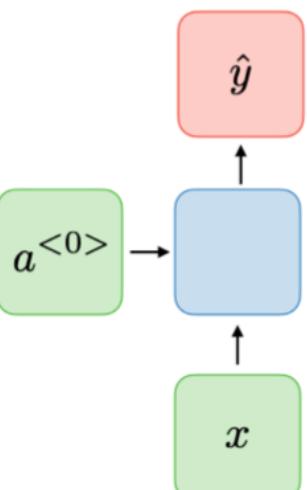
$$a^{<t>} = g_1(W_{aa}a^{<t-1>} + W_{ax}x^{<t>} + b_a)$$

$$y^{<t>} = g_2(W_{ya}a^{<t>} + b_y)$$

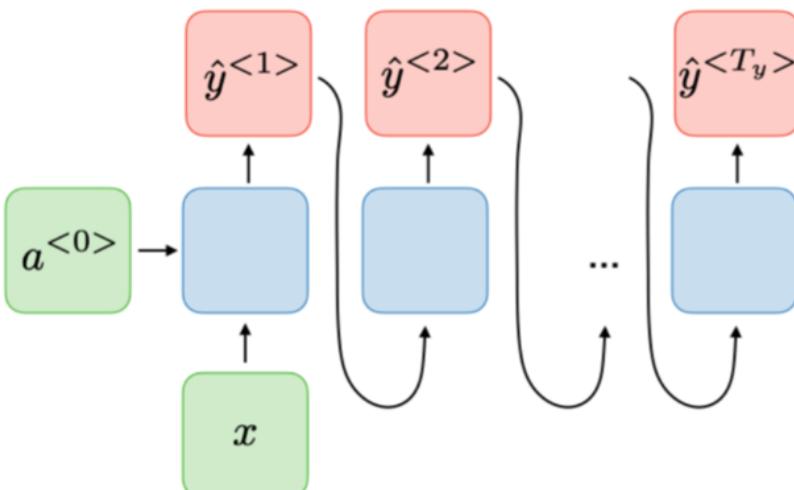
state a at t is a function of x_t and a_{t-1} (and bias)

$W_{ax}, W_{aa}, W_{ya}, b_a$ and b_y are weights that are shared temporally and g_1, g_2 activation functions

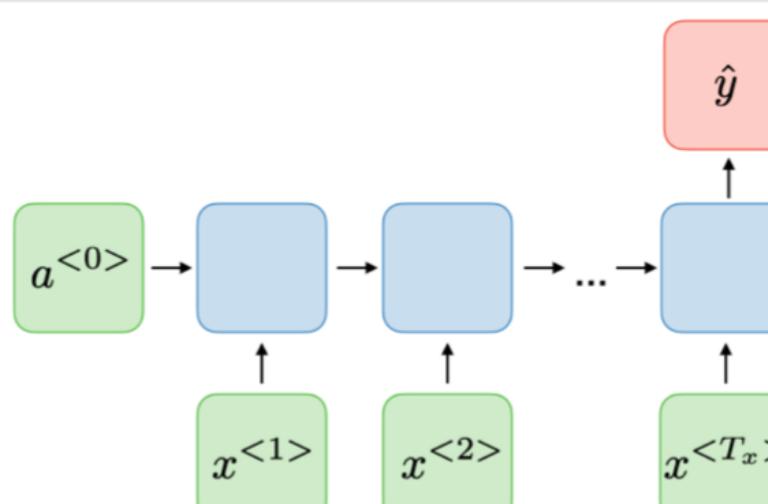
Applications of RNN

Type of RNN	Illustration	Example
One-to-one $T_x = T_y = 1$	 <p>The diagram illustrates a one-to-one Recurrent Neural Network (RNN). It consists of three nodes: an input node labeled x, a hidden state node labeled $a^{<0>}$, and an output node labeled \hat{y}. Arrows indicate the flow of information: an arrow from x to $a^{<0>}$, and another arrow from $a^{<0>}$ to \hat{y}.</p>	Traditional neural network

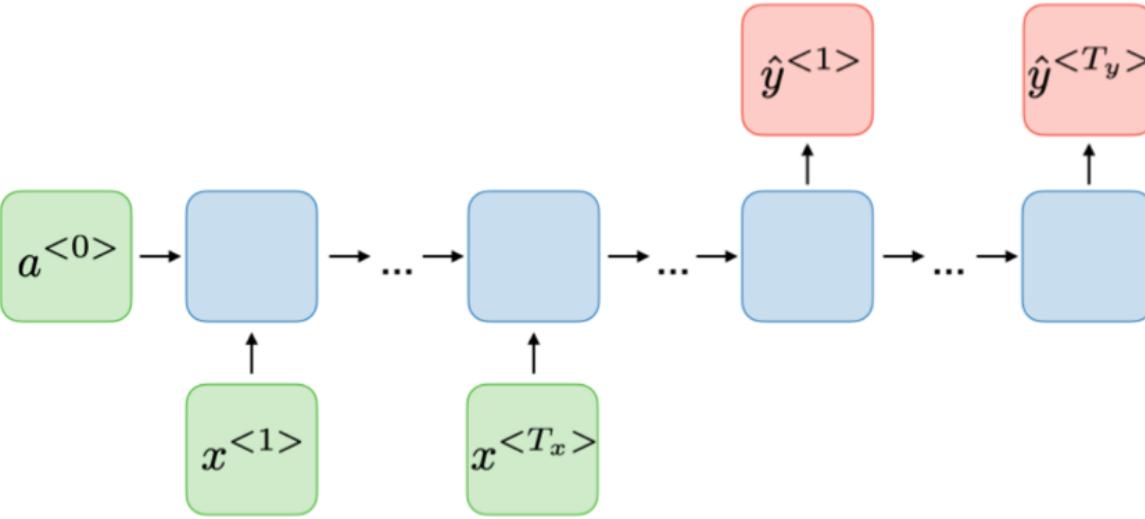
Applications of RNN

Type of RNN	Illustration	Example
One-to-many $T_x = 1, T_y > 1$	 <p>we feed the output into the next layer as input</p>	Music generation

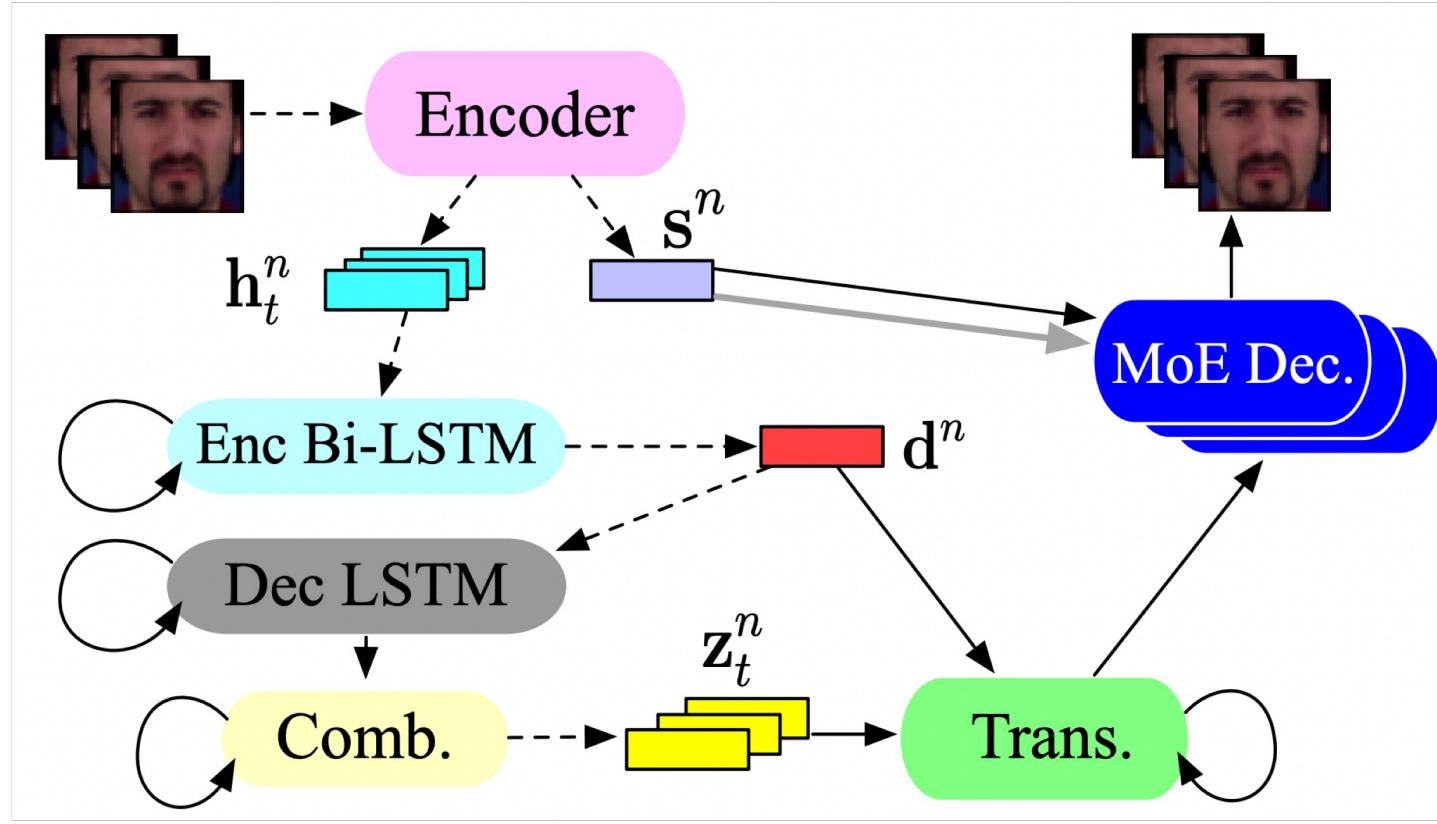
Applications of RNN

Type of RNN	Illustration	Example
Many-to-one $T_x > 1, T_y = 1$	 <p>The diagram illustrates a Many-to-one Recurrent Neural Network (RNN). It starts with an initial hidden state $a^{<0>}$ (green box) which feeds into the first hidden state $a^{<1>}$ (blue box). This process continues sequentially through hidden states $a^{<2>}, \dots, a^{<T_x>}$ (blue boxes). The input sequence consists of tokens $x^{<1>}, x^{<2>}, \dots, x^{<T_x>}$ (green boxes), where each token $x^{<t>}$ is fed into the hidden state $a^{<t>}$. The final hidden state $a^{<T_x>}$ is passed through an output layer (represented by a red box labeled \hat{y}) to produce the output \hat{y}.</p>	Sentiment classification

Applications of RNN

Type of RNN	Illustration	Example
Many-to-many $T_x \neq T_y$	 <p>The diagram illustrates a Many-to-many Recurrent Neural Network (RNN). It consists of two sequences of hidden states. The input sequence (x) starts with a green box labeled $a^{<0>}$ followed by several blue boxes connected by arrows. The output sequence (\hat{y}) starts with a red box labeled $\hat{y}^{<1>}$ followed by several blue boxes connected by arrows. Arrows also point from the input sequence to the first few hidden states, and from the last few hidden states to the output sequence.</p>	Machine translation

VDSM



n

Individual

t

Video Frame

\mathbf{d}^n

Action Performed

\mathbf{s}^n

Identity

\mathbf{z}_t^n

Per-Frame Pose

\mathbf{h}_t^n

Per-image embedding

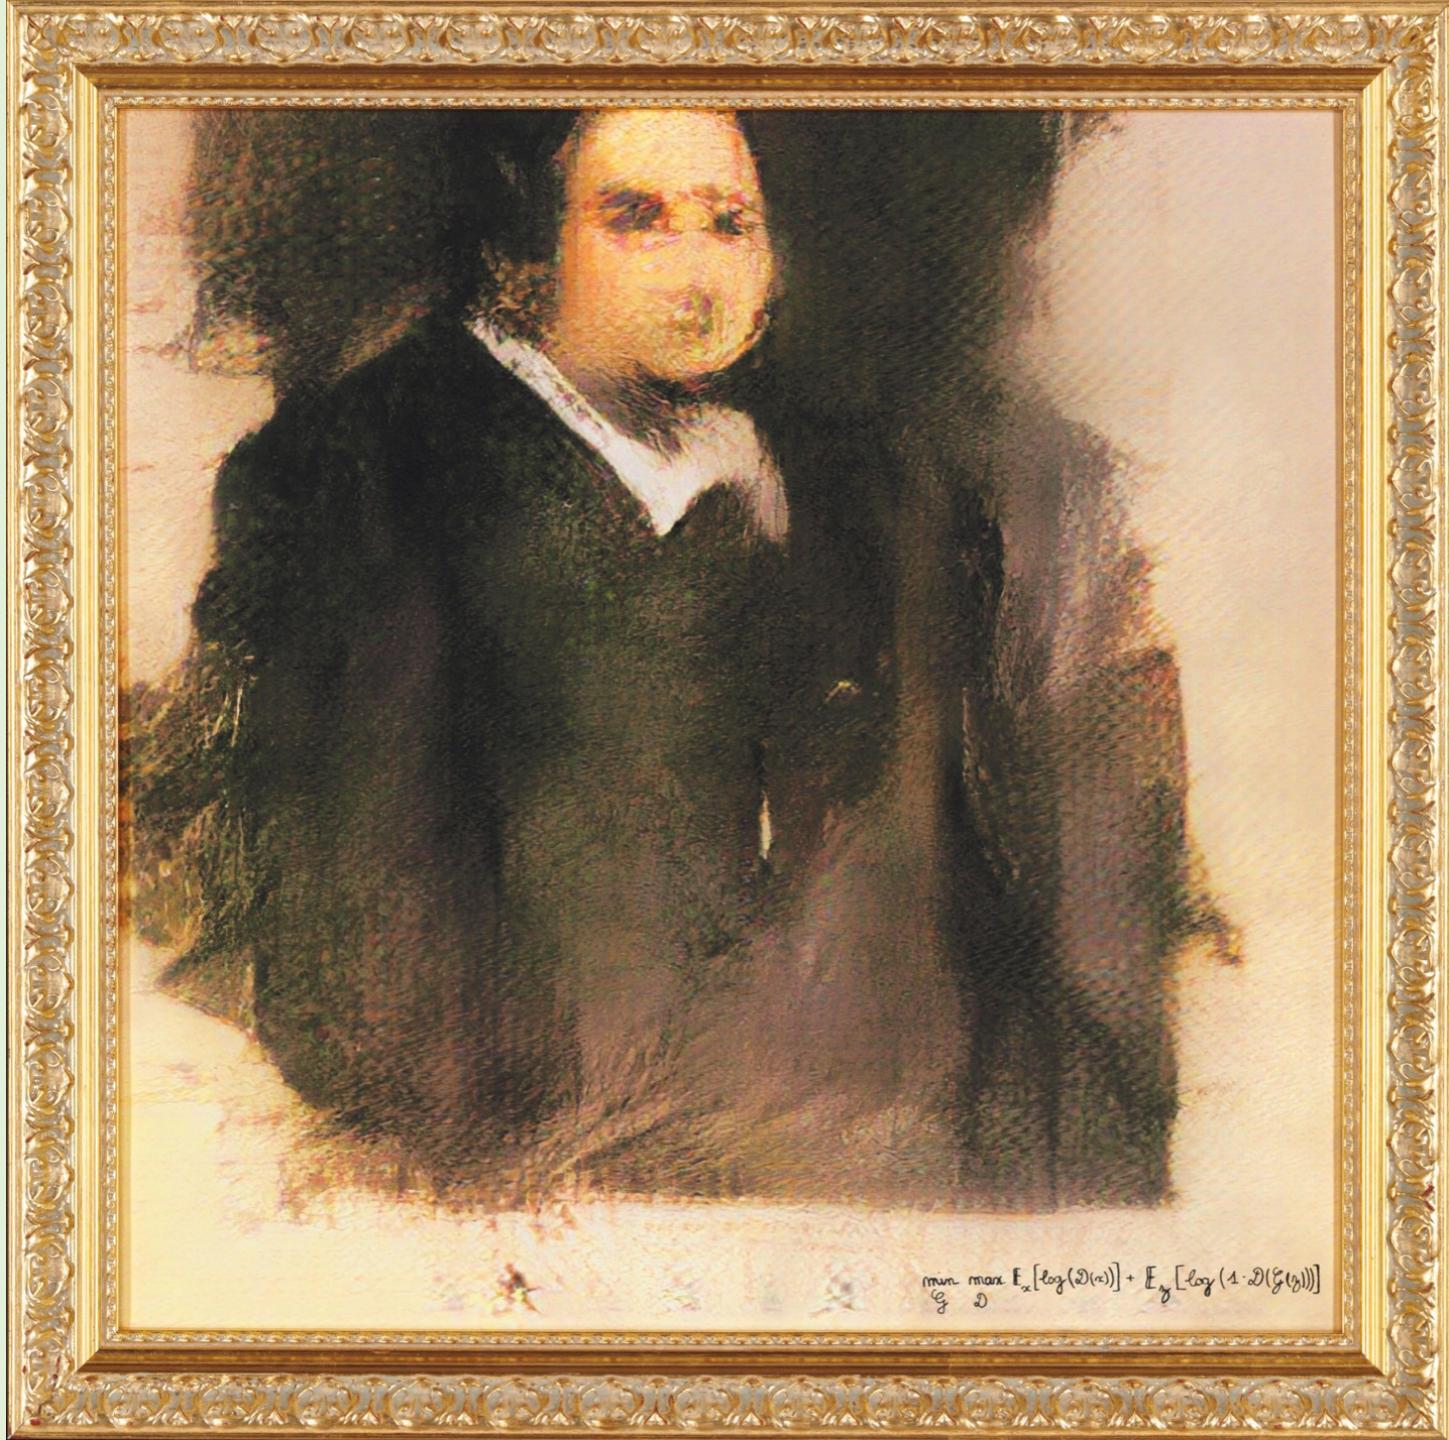
x

Generative Adversarial Network

Invented by Ian Goodfellow

*How much are you
ready to pay for it ?*

Christie's New York 2018



Principle

GENERATOR
“The Artist”
A neural network trying to
create pictures of cats that
look real.



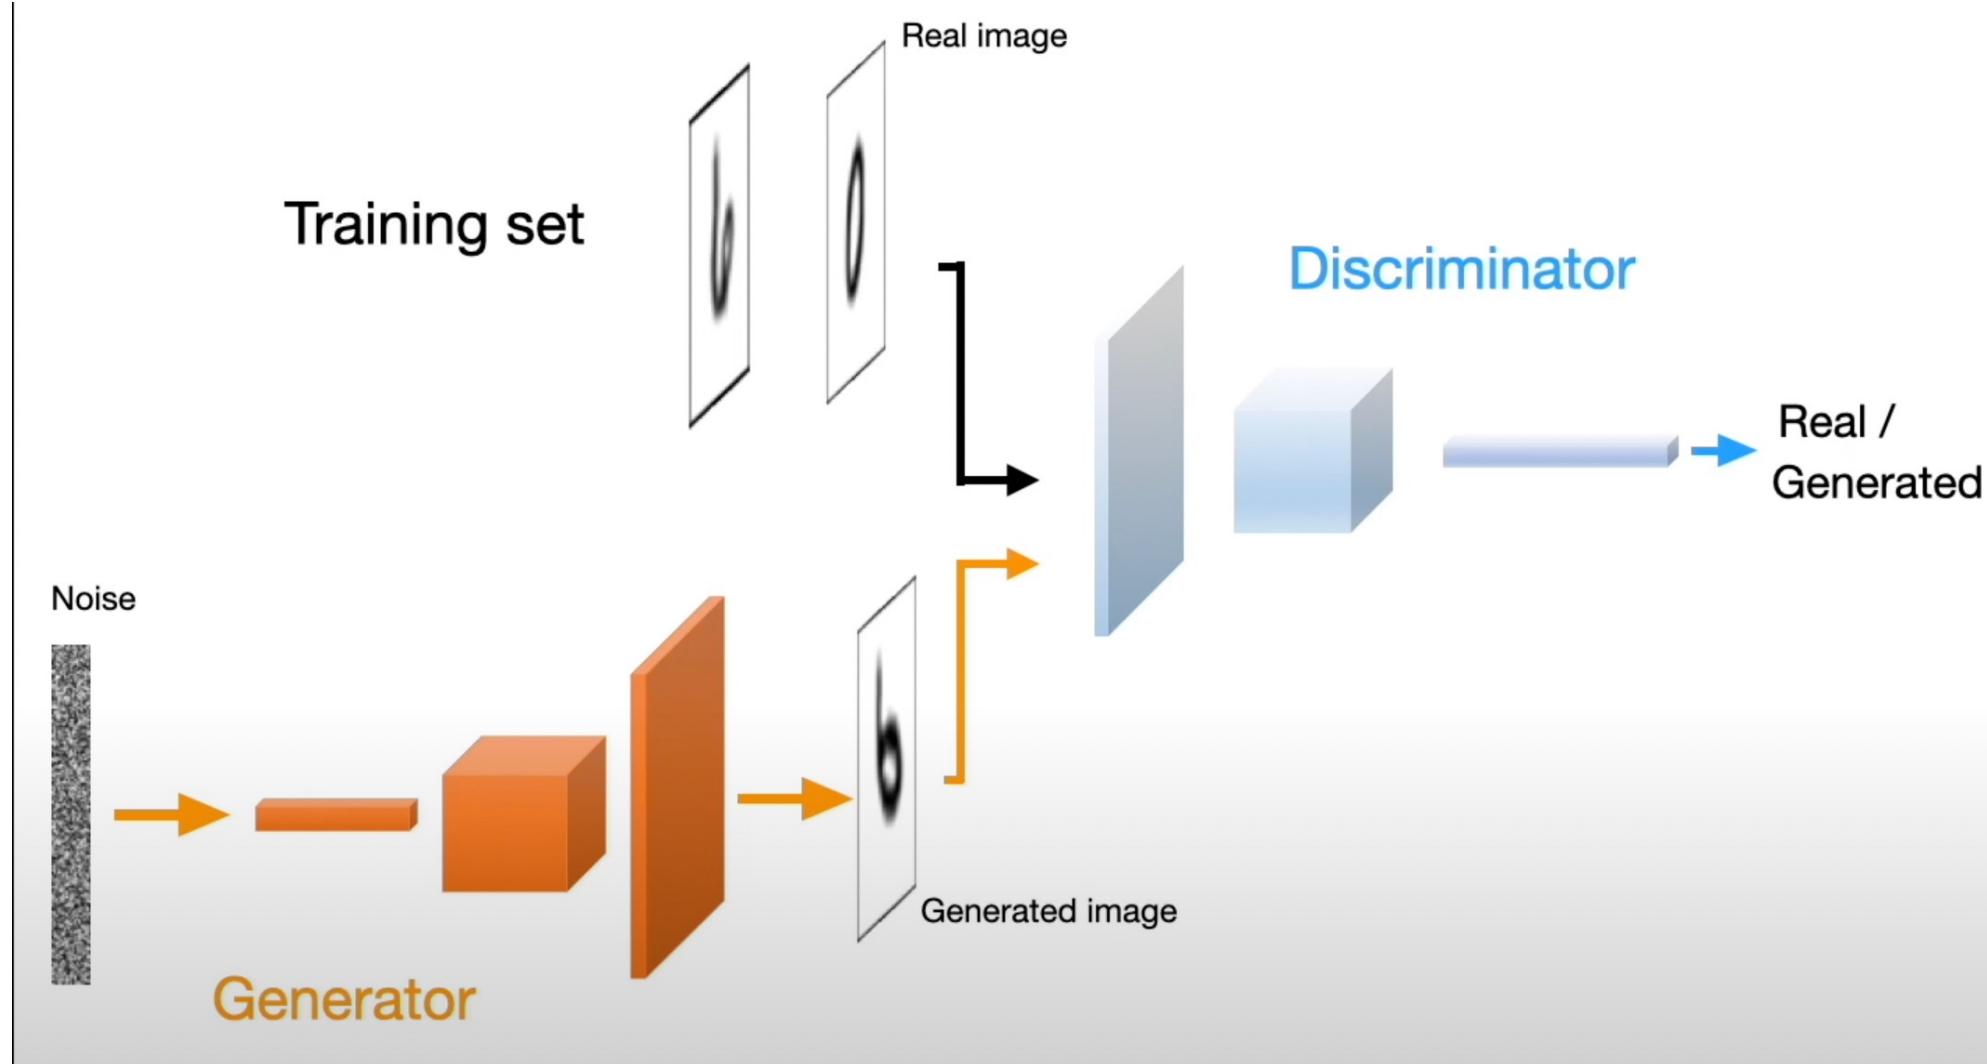
DISCRIMINATOR
“The Art Critic”
A neural network examining
cat pictures to determine if
they’re real or fake.



Thousands of real-world
images labeled “CAT”



Principle



Principle

Step 1: Train Discriminator and ‘Freeze’ Generator parameters

Step 2: Train Generator and ‘Freeze’ Discriminator parameters

two steps in one process, one generator and one discriminator

two optimizers

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}}(\mathbf{x}) [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Principle

Step 1: Train Discriminator and ‘Freeze’ Generator parameters

Step 2: Train Generator and ‘Freeze’ Discriminator parameters

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Discriminator gradient for update (gradient ascent):

predict well on real images
=> want probability close to 1

predict well on fake images
=> want probability close to 0

$$\nabla_{\mathbf{W}_D} \frac{1}{n} \sum_{i=1}^n \left[\overbrace{\log D(\mathbf{x}^{(i)})}^{\text{real images}} + \overbrace{\log (1 - D(G(\mathbf{z}^{(i)})))}^{\text{fake images}} \right]$$

Principle

Step 1: Train Discriminator and ‘Freeze’ Generator parameters

Step 2: Train Generator and ‘Freeze’ Discriminator parameters

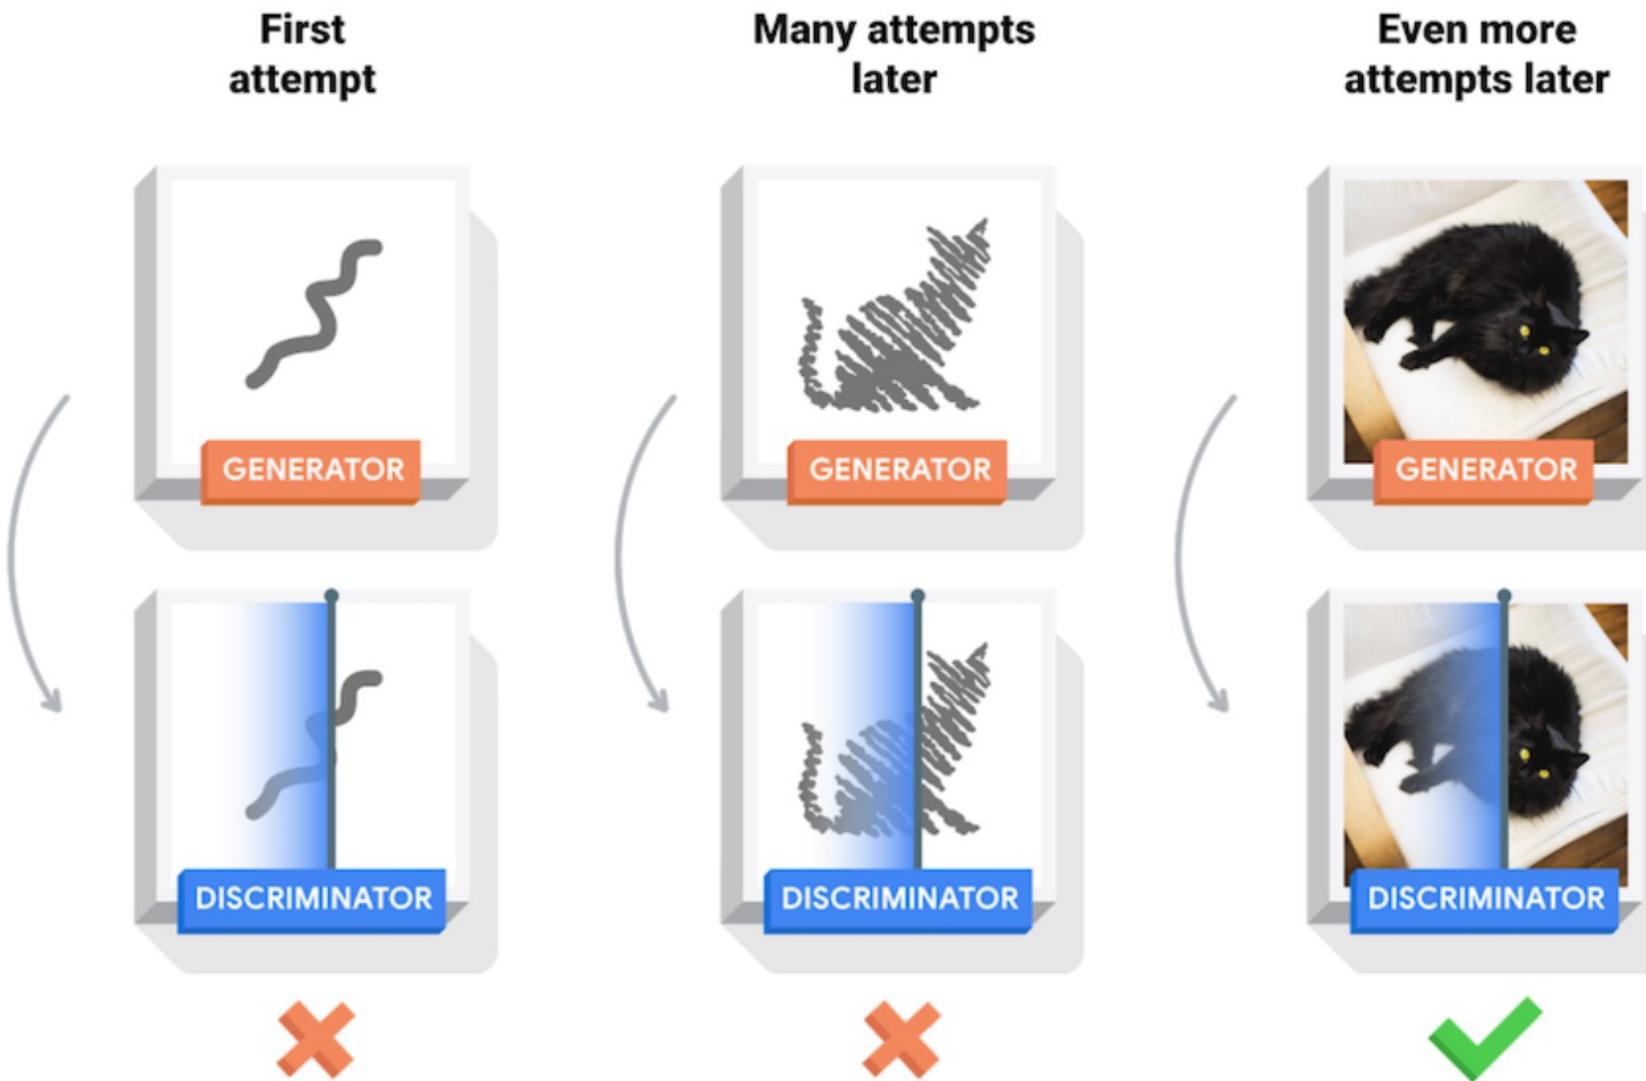
$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_{\mathbf{z}}(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))]$$

Generator gradient for update (gradient descent):

predict badly on fake images
=> want probability close to 1

$$\nabla_{\mathbf{W}_G} \frac{1}{n} \sum_{i=1}^n \log \left(1 - \overbrace{D \left(G \left(\mathbf{z}^{(i)} \right) \right)}^{\text{predict badly on fake images}} \right)$$

Principle



GAN use case : generate images



Figure 3. Example results by our proposed StackGAN, GAWWN [20], and GAN-INT-CLS [22] conditioned on text descriptions from CUB test set. GAWWN and GAN-INT-CLS generate 16 images for each text description, respectively. We select the best one for each of them to compare with our StackGAN.

State Of The Art in GANs



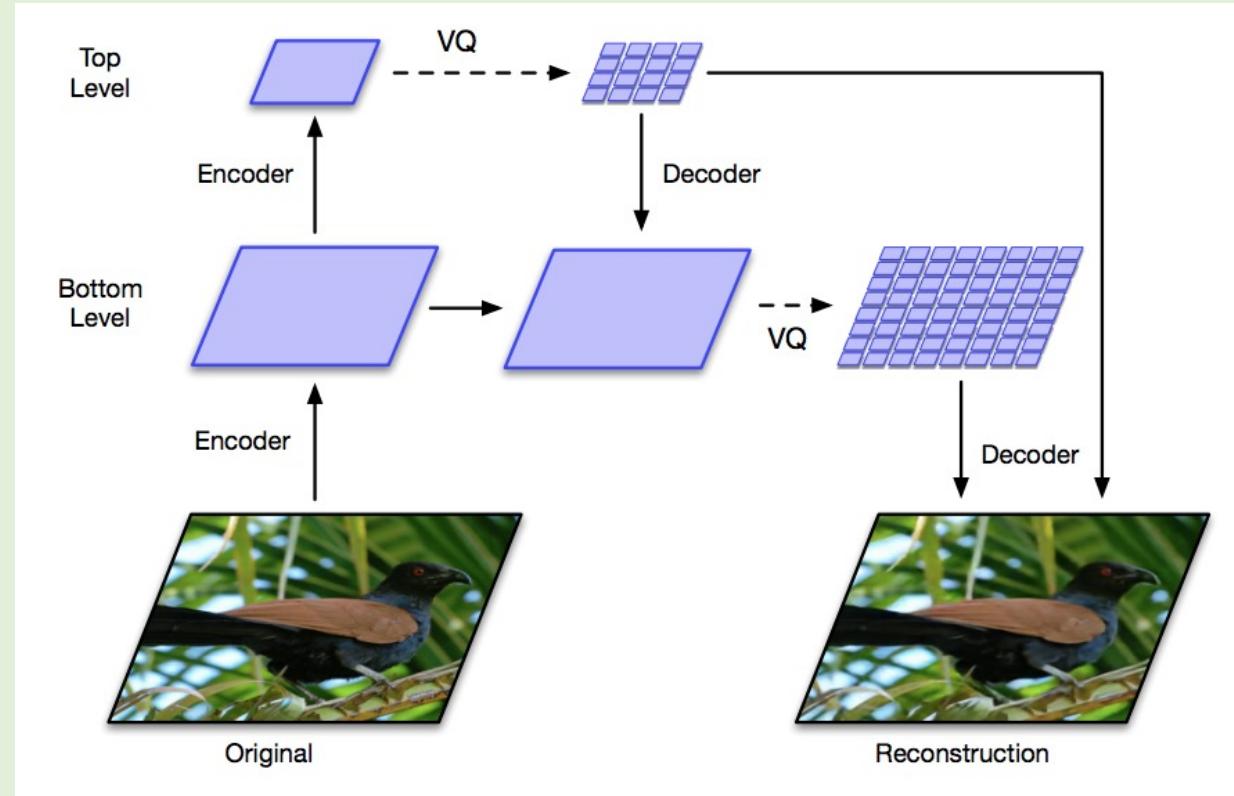
(Karras et al, 2018)

(Brock et al, 2018)

Adversarial Principles are Widely Applicable

- Fairness / Privacy Preservation
- Disentanglement

Variational AutoEncoder



VQ-VAE (2) (van den Oord et al. 2017, Razavi et al. 2019)

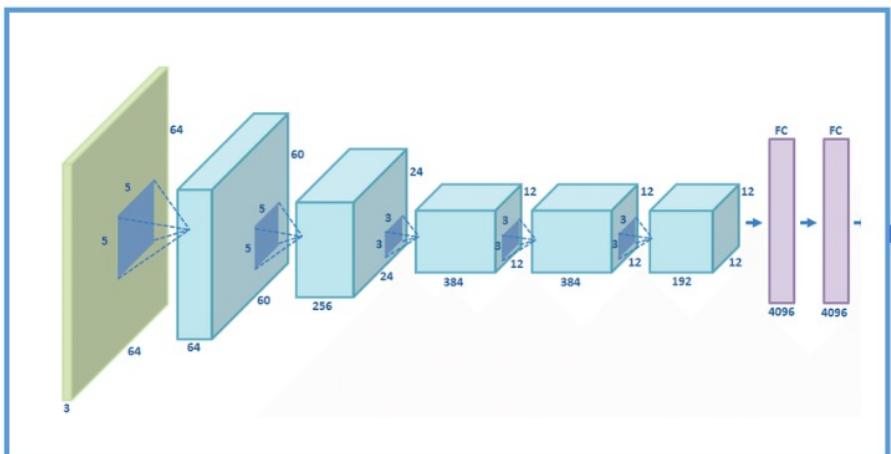
Principle AutoEncoder

Data Generation

- generate appropriately novel data

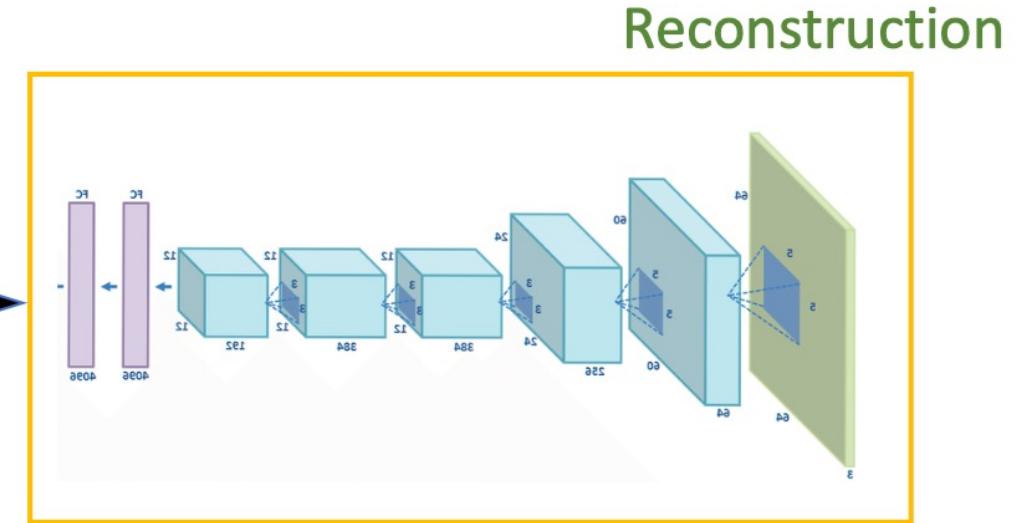


Input



Encoder

Latent space/
Feature

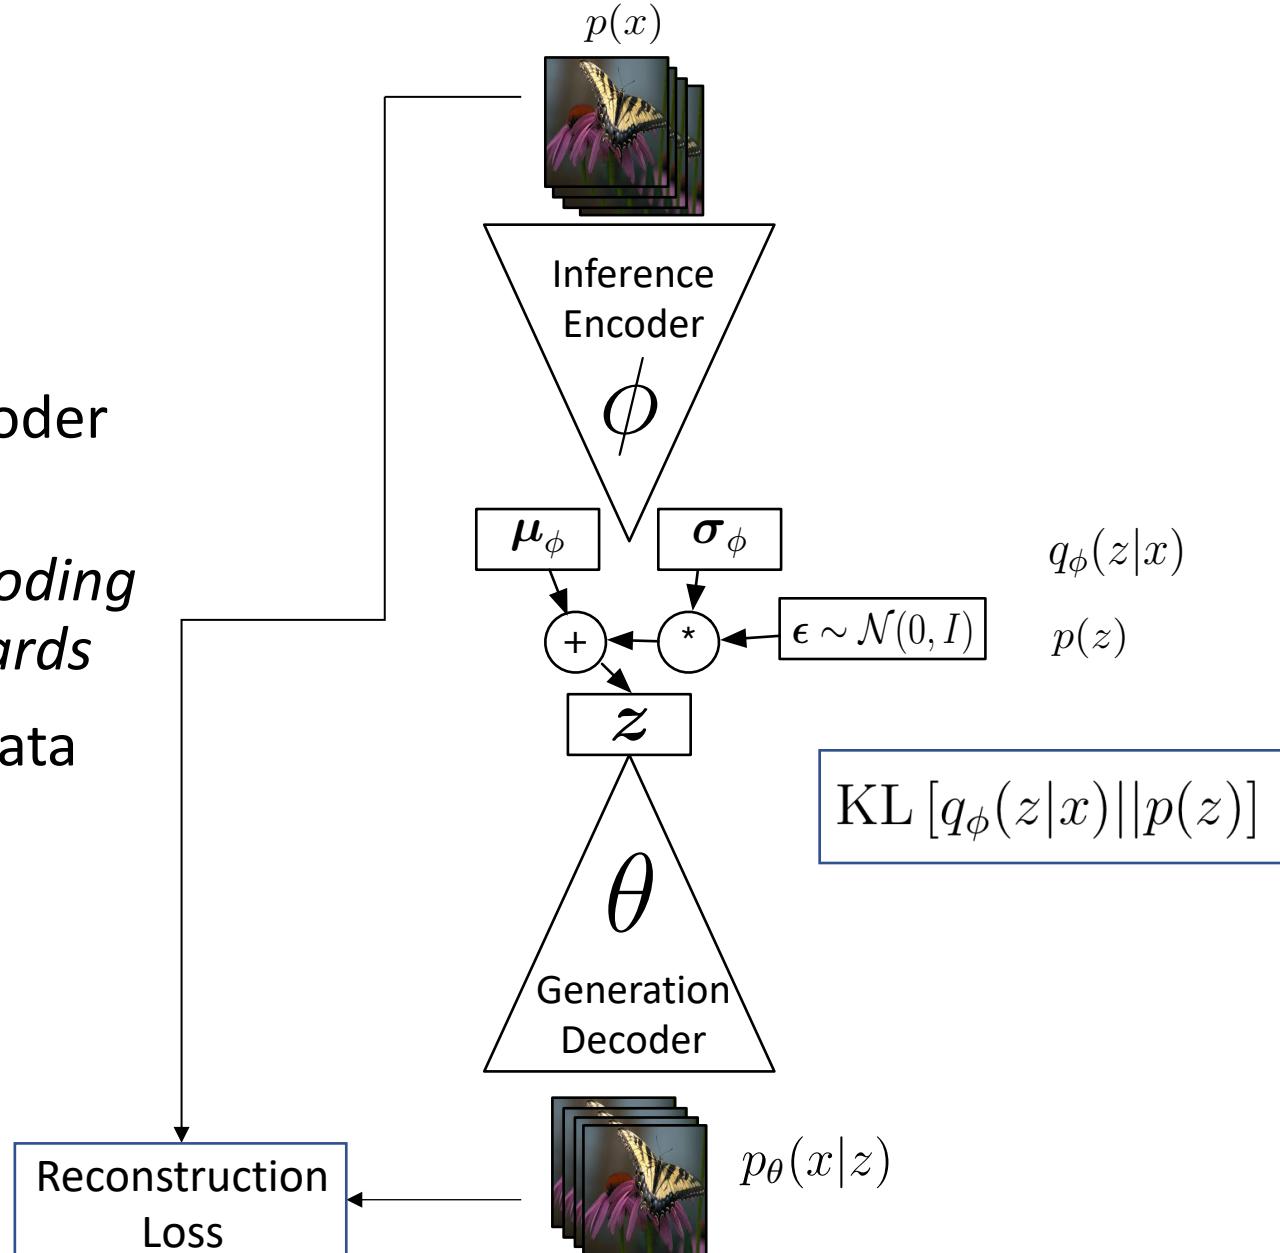


Decoder

Reconstruction

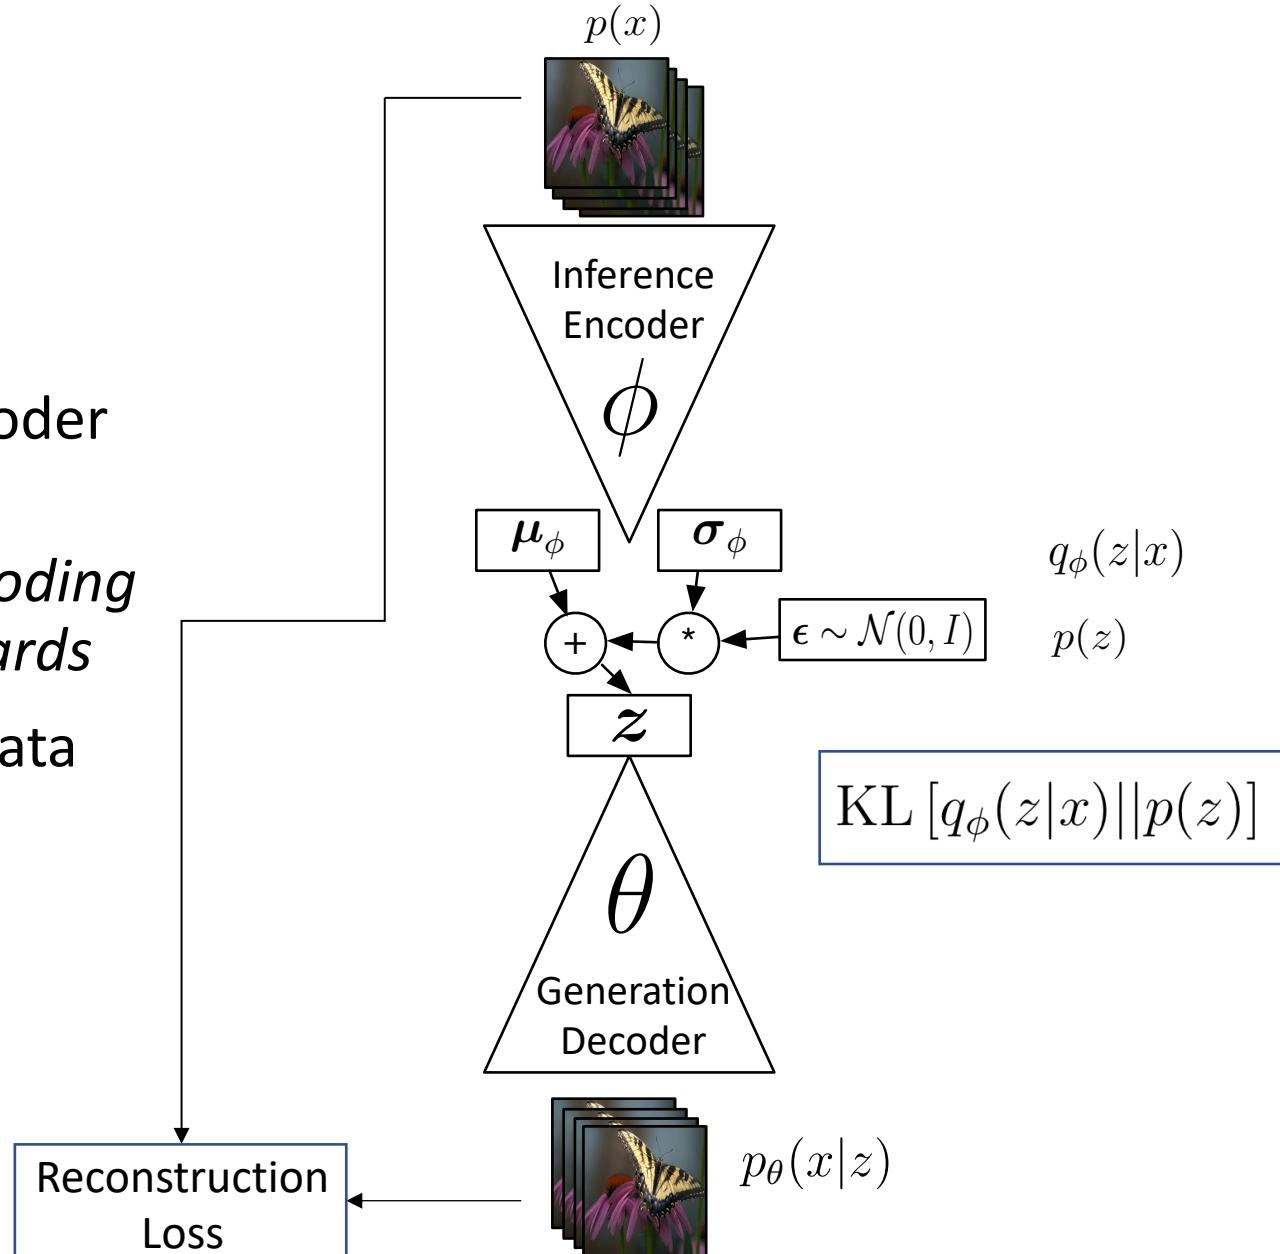
Principle

- Pass images/data in through encoder '*bottleneck*'
- *Parameterize this bottleneck encoding so we can sample from it afterwards*
- Reconstruct the original image/data from the encoding



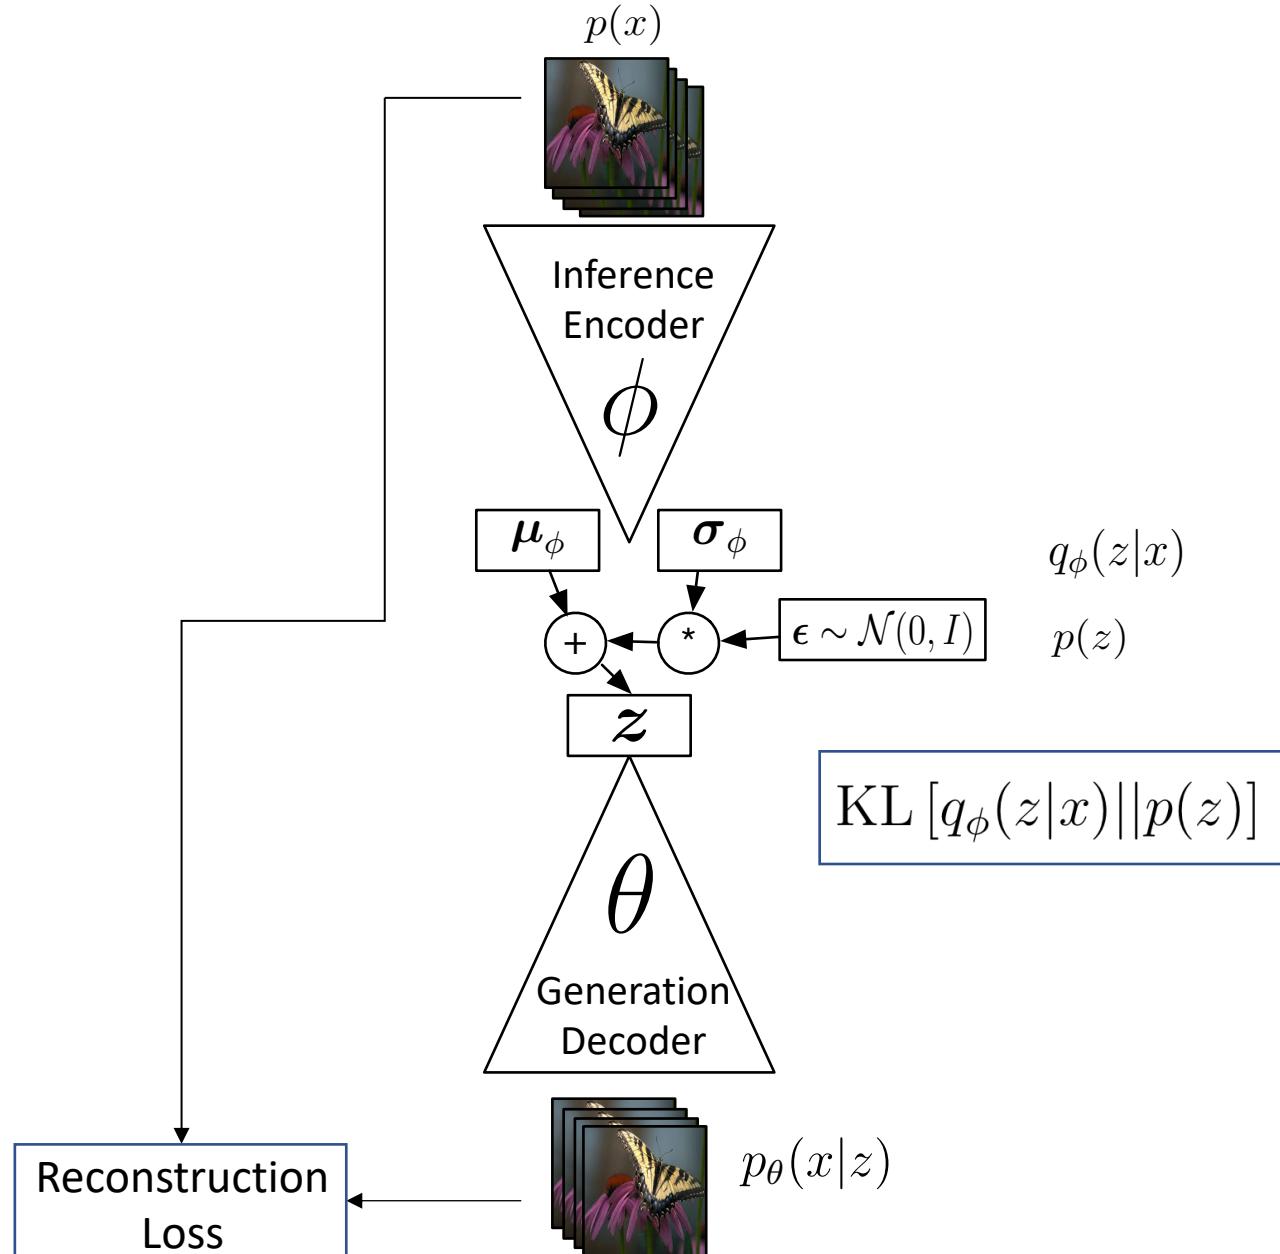
Principle

- Pass images/data in through encoder '*bottleneck*'
- *Parameterize this bottleneck encoding so we can sample from it afterwards*
- Reconstruct the original image/data from the encoding

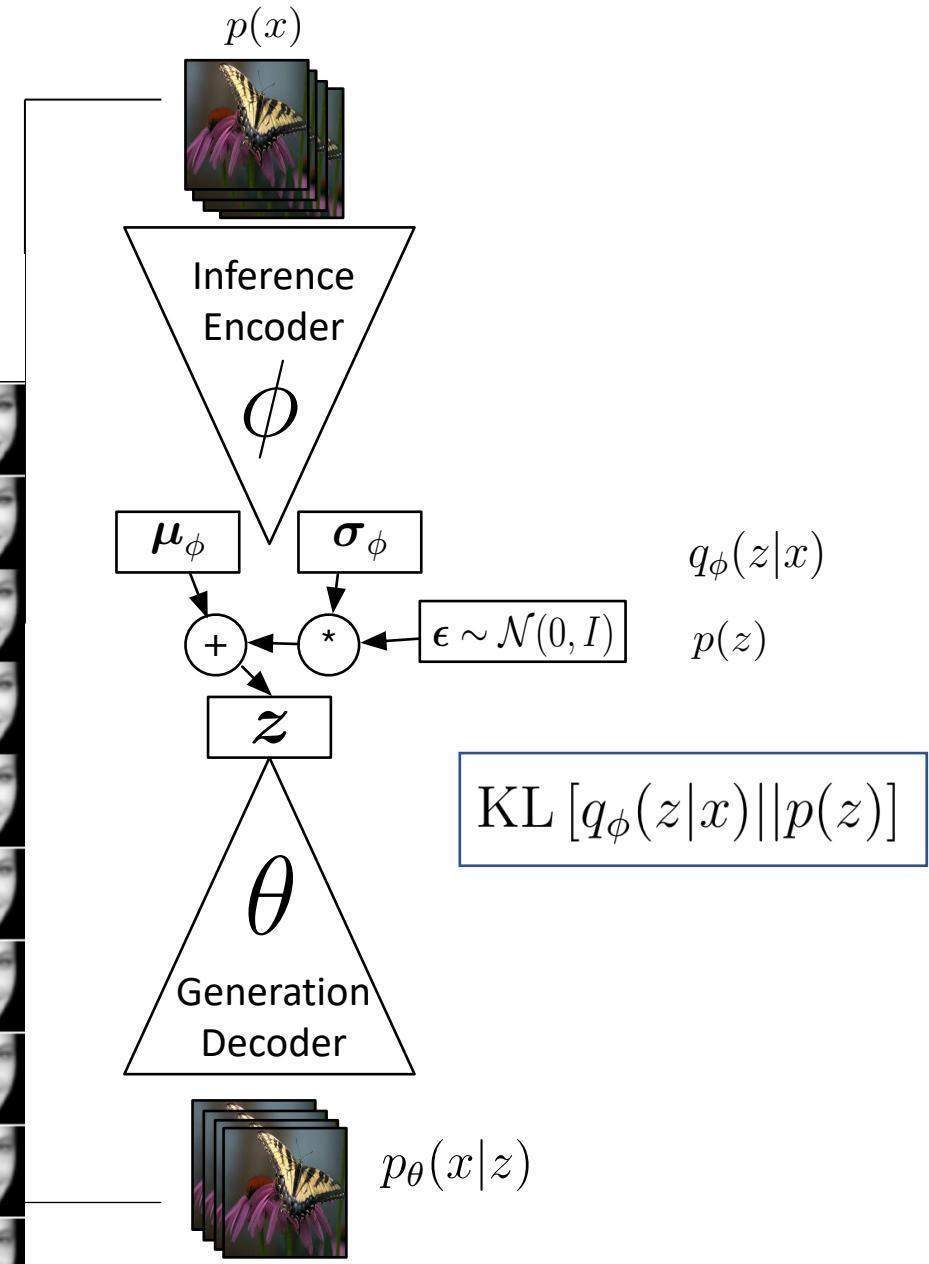
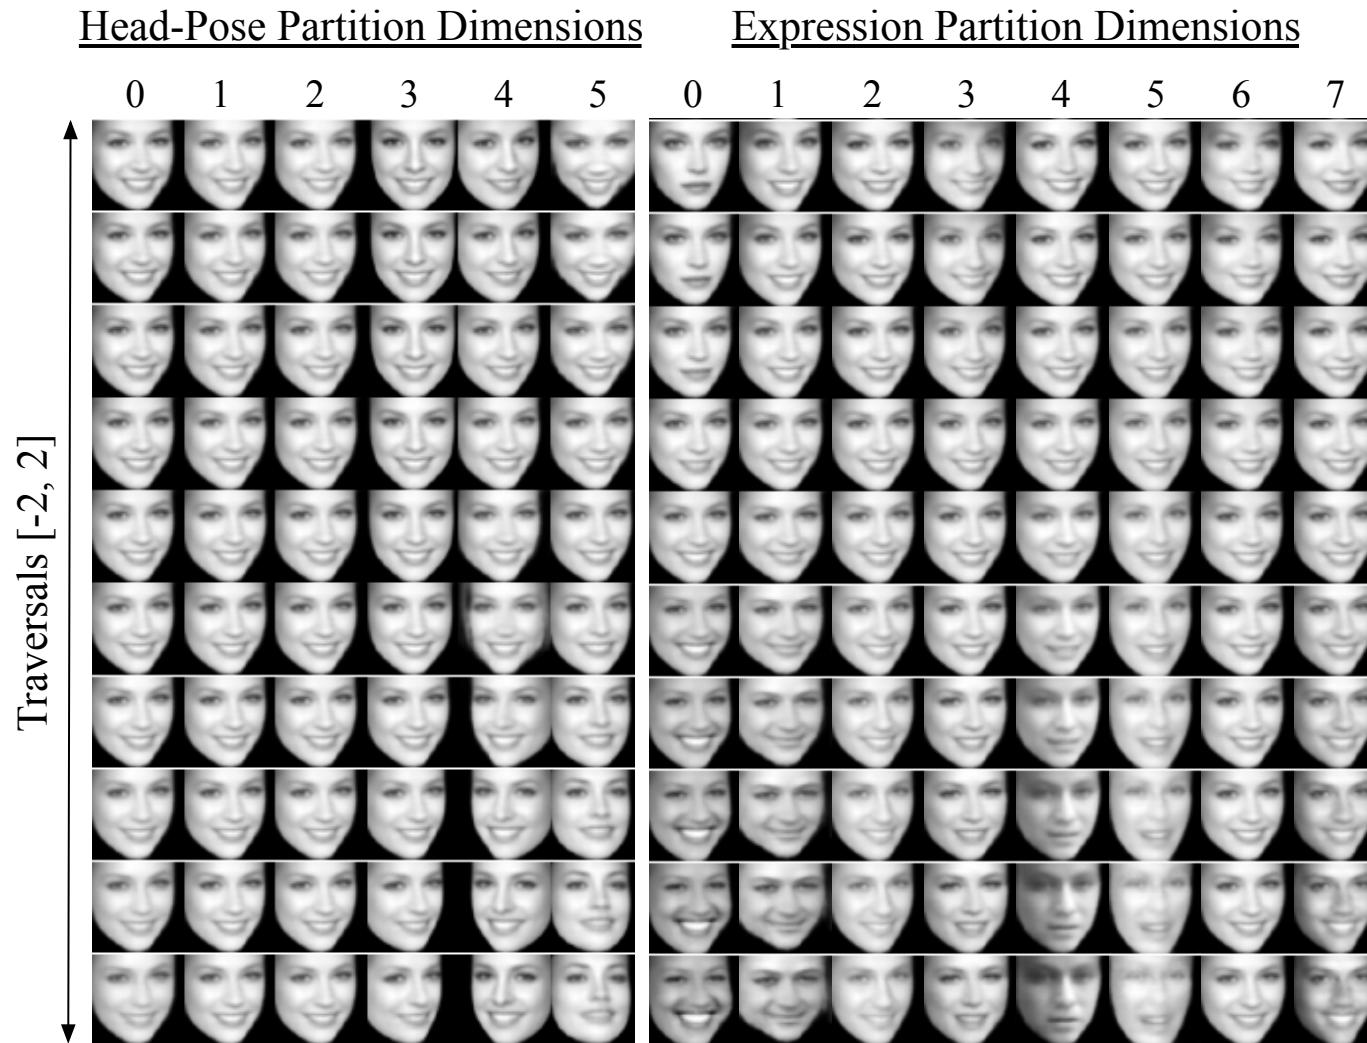


VAE use cases

- Compression
- Data generation
- Latent variable modeling
- Density estimation



VAE use cases



Tutorial / Practical

